Echo State Networks for Time Series Forecasting: Hyperparameter Sweep and Benchmarking

Alexander Häußer[1]

arXiv:2602.03912v5 [cs.LG] 6 Sep 2026

Abstract: This paper investigates the performance of Echo State Networks (ESNs) for univariate forecasting of monthly and quarterly time series from the M4 Forecasting Competition dataset. We evaluate whether a simple first-order autoregressive ESN can serve as a competitive alternative to widely used forecasting methods. The study uses a two-stage design: a *Parameter* dataset is used to analyze ESN model configurations over leakage rate, spectral radius, reservoir size, and regularization selection, while a disjoint *Forecast* dataset is reserved for out-of-sample benchmarking. Forecast accuracy is measured using mean absolute scaled error (MASE), symmetric mean absolute percentage error (sMAPE), and the overall weighted average (OWA). The ESN is compared with simple benchmarks, statistical models including autoregressive integrated moving average (ARIMA), exponential smoothing state space (ETS), the Theta method, and TBATS, as well as multilayer perceptron (MLP), recurrent neural network (RNN), and hybrid exponential smoothing–RNN (ES-RNN) benchmarks. The model-configuration analysis reveals frequency-specific patterns: monthly series tend to favor moderately persistent reservoirs, whereas quarterly series favor more contractive dynamics; across both frequencies, high leakage rates are generally preferred. In the final benchmark, the ESN performs on par with ARIMA and TBATS under mean MASE for monthly data and achieves the second-lowest mean MASE for quarterly data. ES-RNN provides the best aggregate accuracy for both frequencies. Overall, the results indicate that a simple autoregressive ESN can provide competitive forecast accuracy on the considered filtered M4 subsets, particularly under MASE, while requiring low training and forecasting time once the ESN configuration has been fixed.

Keywords: Echo State Networks, Reservoir Computing, Time Series Forecasting, Hyperparameter, M4 Forecasting Competition

1 Introduction

Large-scale forecasting applications often require many time series to be processed automatically under limited computational resources. In such settings, forecasting methods must not only provide accurate predictions, but also be scalable, easy to train, and suitable for heterogeneous time series characteristics. Statistical forecasting methods such as autoregressive integrated moving average (ARIMA), exponential smoothing state space (ETS), the Theta method, and TBATS remain strong benchmarks because they are well established, widely implemented, and often competitive in empirical forecasting studies (Hyndman et al., 2008; Assimakopoulos and Nikolopoulos, 2000; Livera et al., 2011). Fully trained recurrent neural networks (RNNs), by contrast, can capture

[1]Justus-Liebig-University Giessen, Faculty of Economics and Business Studies, Chair of Statistics and Econometrics, Licher Strasse 64, 35394 Giessen, Germany, alexander.haeusser@wirtschaft.uni-giessen.de, 0009-0000-5419-8479

Preprint (September 9, 2026)

nonlinear temporal dependencies but typically require more complex training procedures, larger computational resources, and careful tuning.

Echo State Networks (ESNs), and reservoir computing (RC) more broadly, offer an intermediate approach between these two model classes. ESNs retain recurrent nonlinear dynamics, but only the linear readout layer is trained, while the recurrent reservoir remains fixed after random initialization (Jaeger, 2001; Lukoševičius and Jaeger, 2009; Lukoševičius, 2012). This makes ESNs attractive for automated forecasting settings in which many time series must be modeled with limited manual intervention. The present study therefore examines whether a simple first-order autoregressive ESN can provide competitive forecasting performance on monthly and quarterly time series from the M4 Forecasting Competition dataset (Makridakis et al., 2020).

Although ESNs and RC have been studied in a variety of forecasting contexts, less systematic evidence is available on how simple autoregressive ESNs perform across large heterogeneous benchmark collections of real-world time series under controlled model-configuration searches and standardized forecast-evaluation procedures. Existing work often focuses on specific applications, selected benchmark problems, alternative reservoir topologies, or more complex ESN and RC variants (Viehweg et al., 2025), while providing limited guidance on how basic ESN design choices such as leakage rate, spectral radius, and reservoir size affect forecasting performance across different data frequencies and time series lengths. This issue is particularly relevant when only relatively short time series are available, as is common, for instance, in business forecasting settings where structural breaks, organizational changes, or new product introductions can reduce the usefulness of long historical records.

Time series forecasting competitions such as M3 (Makridakis and Hibon, 2000) and M4 (Makridakis et al., 2020) have highlighted the importance of evaluating forecasting methods on extensive, diverse, and realistic datasets using standardized metrics and benchmark procedures. The M4 dataset, in particular, comprises 100,000 series from multiple application domains and frequencies, providing a stringent testbed for assessing forecast accuracy across heterogeneous time series. While machine learning (ML) and hybrid methods featured prominently among the top-performing approaches in the M4 competition, ESNs have not yet been systematically investigated in this context. It therefore remains an open question whether a comparatively simple, autoregressive ESN can achieve competitive accuracy relative to well-established statistical benchmarks when applied to monthly and quarterly time series with realistic history lengths.

This paper addresses this gap by proposing and empirically evaluating an ESN framework for univariate time series forecasting on a subset of the M4 dataset. The focus is on monthly and quarterly time series, which are particularly important for business and economic forecasting, and on time series with at most 20 years of historical data to reflect practical constraints encountered in many real-world applications. The ESN considered here is purely autoregressive, using only the first lag of the target series as input, and is embedded in a generic preprocessing pipeline that includes stationarity testing using the Kwiatkowski–Phillips–Schmidt–Shin (KPSS) test (Kwiatkowski et al., 1992), differencing where required, and scaling to a symmetric interval. Forecasts are produced recursively over the standard horizons of the M4 Forecasting Competition (18 months and 8 quarters), allowing for a direct comparison with established benchmark methods.

A key contribution of this study is a comprehensive empirical analysis of ESN model configurations at scale. The objective is not to introduce a new ESN architecture or a new training algorithm, but to provide a controlled assessment of how a first-order autoregressive ESN behaves in large-scale univariate time series forecasting. We construct two disjoint random subsets of the M4 data: a *Parameter* dataset used for systematic analysis of ESN hyperparameters and model configurations, and a *Forecast* dataset reserved for out-of-sample evaluation. On the *Parameter* dataset,

we perform an extensive grid search over four ESN design dimensions, i.e., the leakage rate, spectral radius, reservoir size, and choice of information criterion used to select the ridge penalty. This two-stage procedure allows us to examine how ESN design choices affect forecasting accuracy across thousands of heterogeneous time series and to identify frequency-specific configurations that are subsequently evaluated on unseen series.

The predictive performance of the selected ESN configurations is then evaluated on the *Forecast* dataset and compared against a set of widely used benchmark methods. These include simple approaches (naive, Naive2, drift, seasonal naive, and mean forecasts), statistical models (ARIMA, ETS, Theta, and TBATS), and neural network (NN) benchmarks, including the multilayer perceptron (MLP), RNN, and the winning exponential smoothing–recurrent neural network (ES-RNN) model from the M4 Forecasting Competition. Forecast accuracy is assessed using three metrics from the forecasting literature: the symmetric mean absolute percentage error (sMAPE), the mean absolute scaled error (MASE), and the overall weighted average (OWA) (Hyndman and Koehler, 2006; Makridakis et al., 2020). In addition to accuracy, we report computational runtime for the locally estimated methods, providing a joint view of the trade-off between predictive performance and efficiency that is central to large-scale, automated forecasting.

The empirical results indicate that the proposed ESN is competitive with established statistical benchmarks, although its relative performance depends on the data frequency and accuracy metric considered. For monthly series, the selected ESN configuration performs on par with ARIMA and TBATS in terms of mean MASE. For quarterly series, the ESN achieves the second-lowest mean MASE and ranks third according to OWA. ES-RNN provides the best aggregate forecast accuracy for both frequencies, while the standalone MLP and RNN benchmarks perform less favorably. Pairwise statistical tests show that the ESN is generally statistically indistinguishable from the strongest statistical benchmarks. The runtime comparison further shows that the selected ESN configurations can be trained and applied efficiently once the model configuration has been fixed. The analysis of hyperparameters and model configurations reveals systematic frequency-specific patterns: both frequencies favor high leakage rates, but the preferred spectral radius and reservoir size differ between monthly and quarterly data, underlining the importance of adapting reservoir dynamics to data frequency and time series length.

Overall, this study contributes to the empirical forecasting literature in three ways:

1. It evaluates a simple first-order autoregressive ESN framework for univariate time series forecasting in an automated, large-scale setting.

2. It provides an empirical analysis of ESN hyperparameters and model configurations on a heterogeneous benchmark dataset, yielding practical guidance for future applications.

3. It compares the selected ESN configurations with simple, statistical, and NN forecasting methods in terms of forecast accuracy and computational runtime, thereby clarifying the conditions under which this simple ESN specification constitutes a competitive empirical alternative.

The remainder of the paper is structured as follows. Section 2 describes the dataset and sampling strategy. Section 3 details the ESN model, preprocessing steps, and training procedure. Section 4 presents the empirical application, including accuracy metrics, the hyperparameter sweep, random initialization analysis, and benchmark comparison. The paper closes with a summary and concluding remarks. The Appendix provides additional empirical results and documents the computational workflow, software environment, replication code, and computing hardware.

2 Data

The dataset used in the empirical application is taken from the well-known M4 Forecasting Competition (Makridakis et al., 2020). The original M4 dataset consists of 100,000 time series, subdivided into six data frequencies (hourly, daily, weekly, monthly, quarterly, and yearly) and six application fields (Micro, Macro, Industry, Finance, Demographic, and Other). The time series were randomly selected from a database called ForeDeCk, compiled at the National Technical University of Athens (NTUA). The database contains 900,000 time series from multiple publicly available sources. The focus of ForeDeCk is business forecasting; therefore, the time series come from domains such as industries, services, tourism, imports and exports, demographics, education, labor and wage, government, households, bonds, stocks, insurance, loans, real estate, transportation, and natural resources and environment (Spiliotis et al., 2020).

In preparation for the M4 Forecasting Competition, the dataset was scaled to avoid negative values and values lower than 10 to circumvent problems when calculating forecast error metrics. Scaling the data by adding a constant value to the time series ensures that their minimum value equals 10. Additionally, any information pointing to the original time series was removed to guarantee the objectivity of the results. Low-volume and intermittent time series data were not considered to avoid methodological difficulties with close-to-zero and zero values. Therefore, the dataset primarily tests the accuracy of different forecasting methods on continuous time series (Makridakis et al., 2020).

The number of time series for each frequency and application field was determined mainly based on the importance for organizations in terms of operational, tactical, and strategic planning and forecasting. For example, creating monthly and quarterly forecasts is more often required in business and economic forecasting than hourly or yearly forecasts. Moreover, micro and financial data are more frequently used than demographic data (Makridakis et al., 2020; Spiliotis et al., 2020).

In order to evaluate the forecast accuracy of the proposed ESN, only monthly and quarterly time series across all application fields were considered. Some of the time series in the M4 dataset are very long, with more than 200 years of history. While some applications, such as macroeconomic forecasting, have access to extensive historical data, other applications, like business planning and forecasting, typically rely on a few years of historical data, if any. For example, changes in business operations, such as structural or organizational changes within a company, can render older data less applicable. This distinction ensures that our approach remains relevant and applicable to a broad range of real-world scenarios, where only shorter time series are available. Therefore, only time series with at most 20 years of history were retained, corresponding to no more than 240 monthly or 80 quarterly observations.

Given this restriction, 28,820 monthly and 10,950 quarterly time series remain from the available 48,000 monthly and 24,000 quarterly series. From this filtered collection, two disjoint subsets are constructed by random sampling. The first subset, referred to as the *Parameter* dataset, contains 2,400 monthly and 1,200 quarterly time series and is used for the systematic analysis of ESN hyperparameters and model configurations, including leakage rate, spectral radius, reservoir size, and the information criterion used for regularization selection. The second subset, referred to as the *Forecast* dataset, contains the same number of time series and is sampled independently after excluding all series used in the *Parameter* dataset. It is reserved for out-of-sample evaluation and benchmarking against standard forecasting methods.

To ensure reproducibility, the random sampling was conducted using a fixed random seed after applying the frequency and maximum-history filters. The identifiers of the selected M4 series are made available with the replication code.

The *Parameter* dataset is used exclusively to evaluate ESN hyperparameter configurations and to identify the best-performing specification separately for monthly and quarterly data based on average forecast accuracy across all series in that dataset. Once selected, these frequency-specific ESN configurations are fixed and carried forward unchanged to the disjoint *Forecast* dataset, where they are evaluated against the benchmark methods. This two-stage approach helps to mitigate overfitting and provides an out-of-sample assessment of the selected ESN configurations. By structuring the analysis in this way, the study establishes a transparent framework for assessing the forecasting potential of ESNs. Based on the M4 Forecasting Competition, the forecast horizon for the monthly data is 1.5 years ($h = 18$), and for the quarterly data, two years ($h = 8$). Table 8 in the Appendix summarizes the absolute and relative frequencies of both sampled datasets according to the application field.

To characterize the sampled datasets and to assess how they differ from the broader M4 collection, an exploratory data analysis is conducted for key time series characteristics. In particular, we compare the distributions of series length, seasonal strength, and trend strength, since these characteristics are important for many forecasting methods. Some methods explicitly model level, trend, and seasonality, whereas others capture only a subset of these components. The number of observations is also relevant because more flexible, data-driven forecasting methods typically require more historical data than simple benchmark methods.

Figure 1 shows the distributions of the number of observations, the seasonality strength, and the trend strength for the monthly and quarterly time series in both sample datasets and the total dataset for comparison. Details of the calculation for strength of trend and strength of seasonality are provided in the Appendix. The top row visualizes the monthly and the bottom row the quarterly data. The left panel shows the number of observations (i.e., the time series length), the middle panel shows the strength of the seasonality and the right panel shows the strength of trend. The total dataset is colored in black, while the dataset *Parameter* is plotted in orange and *Forecast* is shown in blue. For visual clarity, the x-axis in Figure 1 (left panel) is truncated at 600 observations for the monthly data and 300 observations for the quarterly data when plotting the full M4 reference distributions. This affects only the graphical display: the sampled *Parameter* and *Forecast* datasets remain restricted to series with at most 240 monthly or 80 quarterly observations, and all descriptive analyses are based on the complete filtered samples.

For the monthly series, the full M4 dataset shows an average series length of about 216 observations, with a broad range from 42 to nearly 2,800 observations, reflecting the heterogeneity of the underlying time series sources. The two sampled subsets, however, are shorter on average, with mean lengths of around 128 observations and a maximum of 240, reflecting the imposed maximum-history filter and resulting in shorter, more homogeneous datasets suitable for model experimentation. In terms of seasonality, the average seasonal strength in the total dataset (mean 0.415) is somewhat higher than in the sampled subsets (means around 0.34), suggesting that the selected series may exhibit slightly weaker seasonal components on average. The trend strength remains consistently high across all datasets, with mean values around 0.85–0.89 and median values close to 0.95, indicating that most monthly time series in the M4 dataset display a pronounced trend. The small differences between the total and sampled datasets confirm that the sampling preserved the essential trend characteristics of the original collection.

For the quarterly series, a similar pattern emerges. The full M4 dataset contains an average of about 92 observations per series, while the sampled subsets have a shorter mean length of approximately 56 observations, again reflecting the imposed maximum-history filter. The distribution of seasonal strength shows comparable behavior across datasets, with the sampled subsets having slightly lower mean values (around 0.30–0.31) than the total dataset (0.34), indicating a minor reduction in seasonality. The trend strength is consistently very high across all datasets, with mean

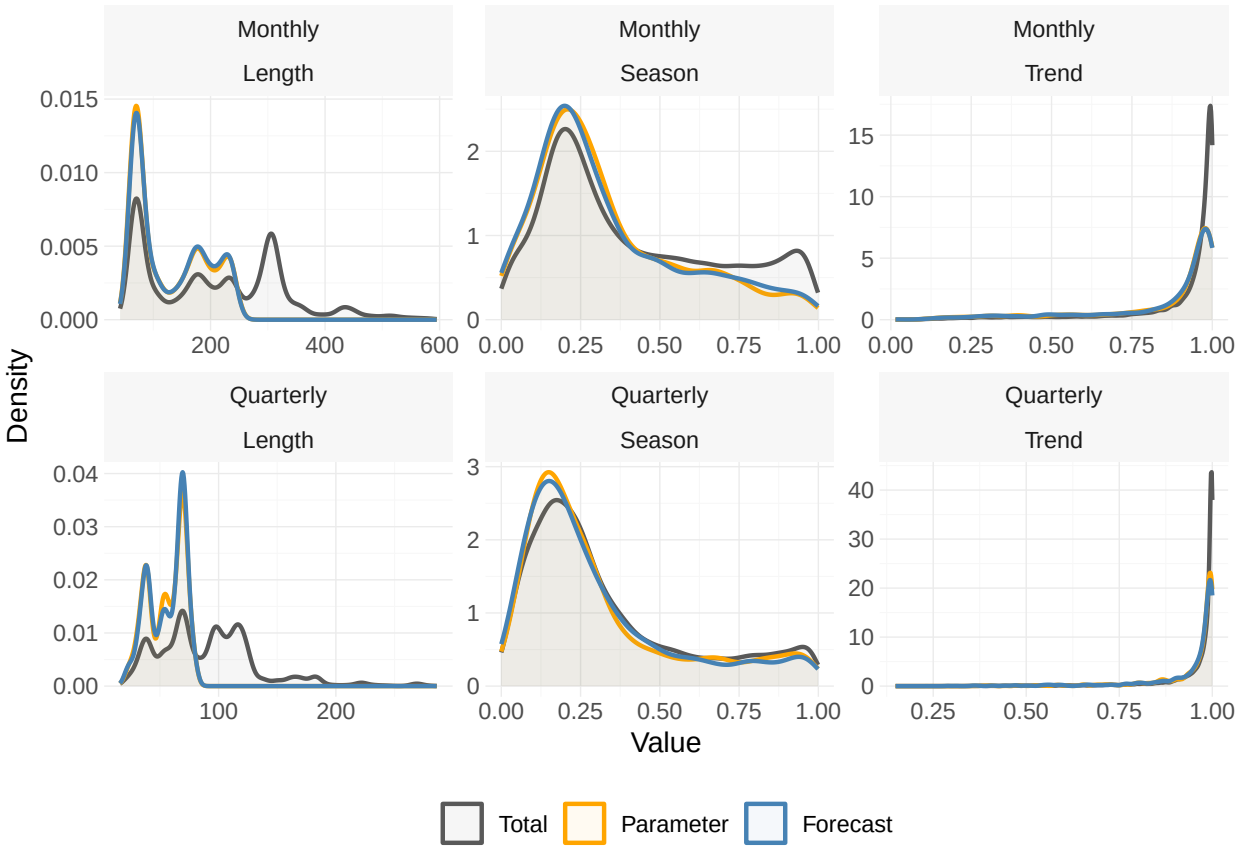

Figure 1: Density plots illustrating the distributions of time series characteristics for the monthly (top row) and quarterly (bottom row) data. The left panels show the number of observations (series length), the middle panels display the strength of seasonality, and the right panels depict the strength of trend. The black lines represent the total M4 dataset, while the orange and blue lines correspond to the randomly sampled *Parameter* and *Forecast* datasets (Data source: M4 Forecasting Competition).

values exceeding 0.94 and median values above 0.98, confirming that most quarterly series in the M4 dataset exhibit strong trending behavior. Tables 9 and 10 in the Appendix provide a numeric summary of the time series characteristics. Figure 2 shows four exemplary monthly time series with varying characteristics.

The sampled datasets preserve several important characteristics of the filtered monthly and quarterly M4 series, particularly the broad distributions of trend strength and series length within the imposed maximum-history constraint. At the same time, the filtering procedure changes the composition of the data: the sampled series are shorter by construction and exhibit slightly weaker seasonality on average than the full monthly and quarterly M4 collections. Consequently, the empirical conclusions should be interpreted as applying to this filtered forecasting setting rather than to the full M4 benchmark.

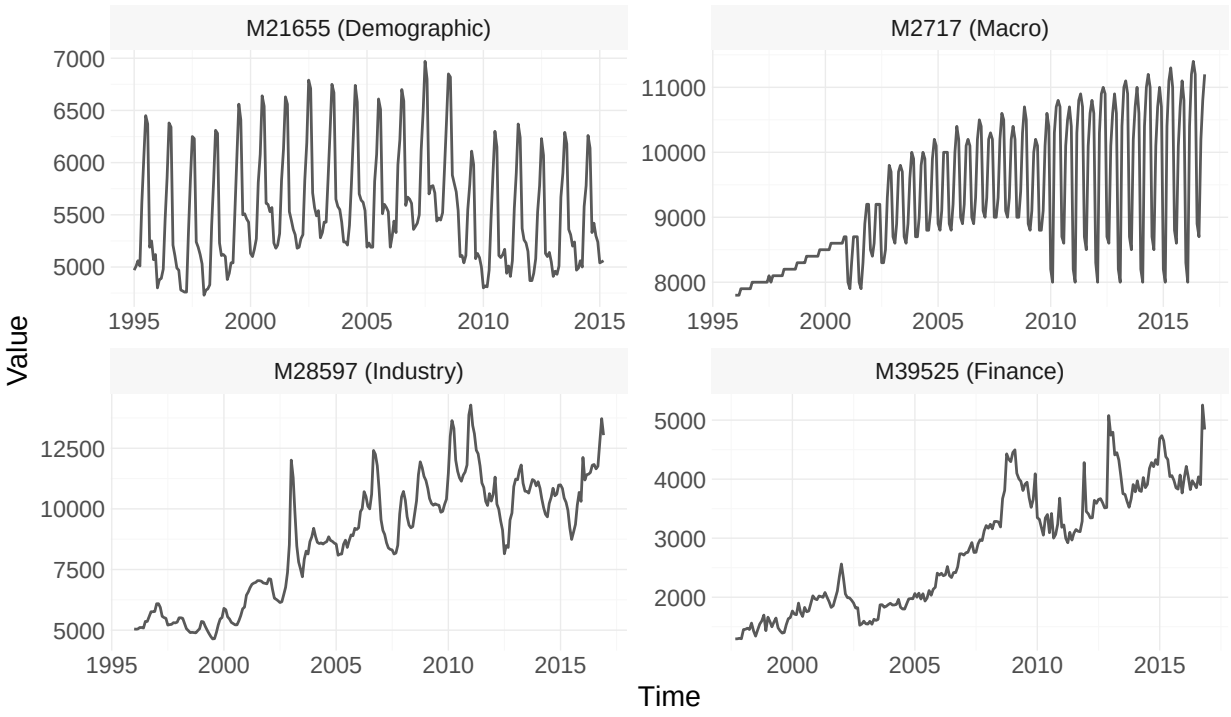

Figure 2: Exemplary monthly time series with varying characteristics. The four time series are taken from the randomly sampled *Parameter* dataset (Data source: M4 Forecasting Competition).

3 Echo State Networks

Forecasting time series in our setting requires learning the relationship between past observations and future values. Since many time series exhibit nonlinear dynamics, purely linear models are often too restrictive to capture these dependencies adequately. A common approach in this context is to map the input into a high-dimensional nonlinear feature space and then estimate the forecast by means of a simple linear readout, such as regression (Lukoševičius and Jaeger, 2009).

ESNs share the general architecture of neural networks but differ fundamentally in training. ESNs separate the nonlinear temporal processing, handled by a fixed, recurrent reservoir, from the readout layer, which is typically linear and trained via supervised learning. This separation contrasts with traditional RNNs, where all weights are jointly optimized (Lukoševičius and Jaeger, 2009).

An ESN consists of an input layer, a large randomly initialized reservoir, and an output layer. Only the output weights $\mathbf{W}^{\text{out}}$ are trained, whereas the input weight matrix $\mathbf{W}^{\text{in}}$ and the reservoir weight matrix $\mathbf{W}$ remain fixed (Lukoševičius, 2012). The reservoir's dynamically evolving internal states provide a rich basis from which the output is linearly reconstructed.

For univariate forecasting, the output layer contains a single output variable, y_t. In the autoregressive ESN considered here, the input is simply the previous observation, $u_t = y_{t-1}$, analogous to a first-order autoregressive model. Owing to this autoregressive setup, ESNs naturally support recursive forecasting ("closed-loop"): the one-step ahead forecast $\hat{y}_{T+1|T}$ is fed back as input to generate the two-step ahead forecast $\hat{y}_{T+2|T}$, and so on. The process continues until the desired forecast horizon h is reached.

This input specification is intentionally restrictive. It keeps the ESN architecture comparable

across all time series and isolates the effect of the reservoir dynamics, but it does not explicitly provide seasonal lagged inputs such as y_{t-12} for monthly data or y_{t-4} for quarterly data. Seasonal information can therefore only enter indirectly through the reservoir state and the recursive forecast dynamics. Consequently, the results should be interpreted as the performance of a simple first-order autoregressive ESN rather than as the best achievable performance of ESNs with richer input structures.

As with most neural networks, input data must be preprocessed before training and forecasting. Many series in our dataset are nonstationary due to trend, seasonality, or heteroscedasticity, and trending behavior is particularly problematic because the reservoir cannot effectively handle such nonstationarity. We therefore apply a Kwiatkowski–Phillips–Schmidt–Shin (KPSS) test (Kwiatkowski et al., 1992) to the training portion of each series to assess whether the output variable y_t is level-stationary. The KPSS null hypothesis is stationarity around a constant level. In the empirical application, the test is evaluated at the 5% significance level. If the null hypothesis of level stationarity is rejected, a single first difference is applied to stabilize the mean and reduce nonseasonal trend components (Hyndman and Athanasopoulos, 2021). Seasonal differencing is not used in the baseline specification.

After detrending, the data are scaled to the interval $[-0.5, 0.5]$. Scaling is important because ESNs are sensitive to the magnitude of the inputs, and keeping the data within a controlled range improves numerical stability and reservoir dynamics. Empirically, scaling to the symmetric range $[-0.5, 0.5]$ performed well, though other normalization methods (e.g., min–max or z-score standardization) are also suitable.

Scaling parameters are estimated exclusively from the transformed training data and are then applied during recursive forecasting to avoid information leakage. After forecasts are generated on the transformed scale, the inverse scaling transformation is applied first, followed by reconstruction of the original scale by reversing the differencing operation using the last available training observation.

In an ESN, the hidden layer (reservoir) is represented by $\mathbf{x}_t = (x_{t,1}, x_{t,2}, \ldots, x_{t,N_x})^\top$, where $x_{t,n}$ denotes the internal state of reservoir unit n at time t, and N_x is the total number of reservoir units (reservoir size). Given a single input $u_t = y_{t-1}$, the reservoir dynamics follow the standard update equation

$$\tilde{\mathbf{x}}_t = \tanh\left(\mathbf{W}^{\text{in}} u_t + \mathbf{W}\mathbf{x}_{t-1}\right), \tag{1}$$

for $t = 1, 2, \ldots, T$, with initial condition $\mathbf{x}_0 = \mathbf{0}$. The matrix $\mathbf{W}^{\text{in}} \in \mathbb{R}^{N_x \times 1}$ denotes the input weight matrix, while $\mathbf{W} \in \mathbb{R}^{N_x \times N_x}$ denotes the final reservoir weight matrix used in the state update. The nonlinear transformation in (1) is applied element-wise via the hyperbolic tangent activation function, which is the most common choice in ESN applications (Lukoševičius, 2012).

The input weight matrix $\mathbf{W}^{\text{in}}$ is typically initialized with random values, commonly drawn from a uniform distribution on $[-0.5, 0.5]$, although other distributions such as the normal distribution are also possible. In the present study, input scaling is held fixed through this initialization range, which reduces the dimensionality of the hyperparameters and model configurations. Since input scaling can affect ESN performance, the results should be interpreted conditional on this fixed choice for input scaling. For the reservoir, let $\mathbf{W}_0$ denote the initial random weight matrix. In the present application, $\mathbf{W}_0$ is constructed to be sparse, with a density of 50%, meaning that half of its entries are randomly drawn from a uniform distribution on $[-0.5, 0.5]$ and the remaining half are set to zero. The reservoir density is held fixed at 50% to limit the dimensionality of the already extensive hyperparameter sweep and to isolate the effects of leakage rate, spectral radius, reservoir size, and regularization selection. The value is therefore used as a pragmatic baseline rather

than as a theoretically or empirically optimal density. Since reservoir sparsity can interact with the spectral radius, the echo state property, and the rate of memory decay, the empirical results should be interpreted as conditional on this fixed density. Substantially sparser reservoirs are also common in ESN applications, and a systematic sensitivity analysis of reservoir sparsity remains a direction for future research.

A crucial ESN hyperparameter is the spectral radius of the reservoir matrix. Let $\rho(\mathbf{W}_0)$ denote the spectral radius of the initial reservoir matrix $\mathbf{W}_0$, defined as its largest absolute eigenvalue. We first normalize $\mathbf{W}_0$ according to $\mathbf{W}_1 = \mathbf{W}_0/\rho(\mathbf{W}_0)$ so that $\mathbf{W}_1$ has unit spectral radius. The final reservoir weight matrix is then obtained as $\mathbf{W} = \rho\mathbf{W}_1$, where ρ is the target spectral radius used as a hyperparameter. In this way, ρ controls the effective scale of the recurrent weight matrix and thereby influences the reservoir dynamics.

The spectral radius is commonly used as a heuristic control parameter for the memory and stability of the reservoir. Smaller values usually lead to more contractive dynamics and faster forgetting, whereas larger values allow past inputs to influence the reservoir state for longer. A related concept is the echo state property (ESP), which requires that the influence of past inputs and initial conditions fades over time so that the current reservoir state is primarily driven by the recent input history. The spectral radius can influence this behavior, but the ESP is not determined by the spectral radius alone; it also depends on leakage, input scaling, the activation function, and the realized reservoir weights. Therefore, values close to or above one should be interpreted as exploratory settings near the boundary of stable reservoir dynamics rather than as formal guarantees of the ESP (Lukoševičius, 2012).

In the empirical application, the reservoir size is chosen according to the rule of thumb $N_x = \min(\lfloor \tau T \rfloor, 200)$, with τ as the so-called reservoir scaling parameter and T denoting the sample size. Thus, the number of reservoir units is proportional to the available observations, rounded down using the floor operator. The upper bound of 200 prevents excessively large reservoirs for long time series. Although this rule is linear in T, nonlinear alternatives could also be considered. This rule of thumb is a pragmatic design choice adopted for the empirical application and is not derived from a formal theoretical result. Its purpose is to provide a simple and consistent way to adapt reservoir size to time series length while keeping the overall model complexity computationally manageable across a large number of series.

For a leaky integrator ESN, a leakage rate $\alpha \in (0, 1]$ is introduced:

$$\mathbf{x}_t = (1 - \alpha)\,\mathbf{x}_{t-1} + \alpha\tilde{\mathbf{x}}_t \tag{2}$$

This update blends the previous internal state with the newly computed nonlinear state, introducing a controllable form of memory analogous to exponential smoothing. Setting $\alpha = 1$ recovers the standard (non-leaky) ESN, while $0 < \alpha < 1$ produces a weighted average of $\mathbf{x}_{t-1}$ and the $\tanh(\cdot)$ term. Smaller values of α slow the reservoir dynamics and enable the network to capture longer-term dependencies (Lukoševičius, 2012).

The internal states are computed according to equations 1 and 2 and, together with an intercept term, stacked over time to form the design matrix $\mathbf{X}$. Accordingly, each row of $\mathbf{X}$ is of the form $(1, \mathbf{x}_t^\top)$, where the first column contains ones and the remaining columns contain the reservoir states at time t.

Initializing the reservoir at $\mathbf{x}_0 = 0$ can lead to an artificial transient phase during which the reservoir dynamics have not yet stabilized. It is therefore standard practice to discard an initial portion of the observations, often referred to as a washout or warm-up period, before estimating the output weights (Lukoševičius, 2012). Based on preliminary empirical experimentation conducted as part of the present study, the first $\delta = \lfloor 0.05T \rfloor$ observations are removed from the de-

sign matrix. This proportional rule allows the washout period to scale with the time series length while retaining 95% of the observations for readout estimation. The value should be interpreted as a pragmatic implementation choice rather than as a theoretically optimal washout length.

Overfitting is a common challenge in machine learning and particularly in ESNs, where models may capture both signal and noise. The goal is to obtain a regularized model with good forecasting performance, which requires controlling overfitting and reducing prediction error. In the ESN literature (Jaeger, 2001, 2002), ridge regression is used for training the linear readout. When many (often highly correlated) features are present, linear models can suffer from multicollinearity, which may impair forecast performance; ridge regression mitigates these issues (Hoerl and Kennard, 1970).

Regularization helps stabilize ESN training by addressing overfitting and ill-posed optimization. Ridge regression resembles least squares, but shrinks coefficients toward zero. Given the output vector $\mathbf{y}$ and the design matrix $\mathbf{X}$ of internal states, the linear model is

$$\mathbf{y} = \mathbf{X}\mathbf{W}^{\text{out}} + \boldsymbol{\epsilon}, \tag{3}$$

where $\mathbf{y}$ is $T \times 1$, $\mathbf{X}$ is $T \times (N_x + 1)$ and includes an intercept, and $\boldsymbol{\epsilon}$ is an error term.

Ridge regression introduces a penalty on coefficient size, analogous to weight decay in neural networks (Hastie et al., 2009). The estimator is defined as

$$\widehat{\mathbf{W}}^{\text{out}} = \left(\mathbf{X}^\top \mathbf{X} + \mathbf{R}_\lambda\right)^{-1} \mathbf{X}^\top \mathbf{y}, \tag{4}$$

where $\mathbf{R}_\lambda = \text{diag}(0, \lambda, ..., \lambda)$ ensures that the intercept is unpenalized while all N_x coefficients share a common penalty λ. Adding $\mathbf{R}_\lambda$ resolves rank-deficiency in $\mathbf{X}^\top \mathbf{X}$ and thereby stabilizes the matrix inversion (Hastie et al., 2009).

Selecting an appropriate λ requires a transparent and computationally feasible tuning strategy. In this study, information criteria are used as pragmatic selection rules for the ridge penalty within each fixed ESN configuration. Because the reservoir induces a nonlinear transformation of the original time series, these criteria should not be interpreted as exact likelihood-based model selection criteria for the full ESN process. Rather, they provide a consistent way to compare ridge readouts fitted to the same reservoir-state representation. Model complexity is measured through the effective degrees-of-freedom, defined via the trace of the ridge hat matrix (Hastie et al., 2009):

$$df_\lambda = \text{tr}\left(\mathbf{H}_\lambda\right) = \text{tr}\left(\mathbf{X}\left(\mathbf{X}^\top \mathbf{X} + \mathbf{R}_\lambda\right)^{-1} \mathbf{X}^\top\right). \tag{5}$$

Several hyperparameter search methods exist, but grid search and random search remain widely used (Bergstra and Bengio, 2012). Grid search evaluates all points in a predefined grid, while random search samples values from a bounded domain, often uniformly. Bergstra and Bengio (2012) show that random search is typically more efficient, often finding equal or better models with far fewer evaluations.

This paper employs random search over $\lambda \in [10^{-4}, 2]$, an interval chosen through preliminary experimentation. Specifically, for each fixed ESN configuration defined by α, ρ, and τ, we draw $K = 2N_x$ candidate values of λ uniformly, thereby scaling the search size with the reservoir dimension. For each candidate, we fit the linear readout and compute the corresponding information criterion, which is used solely to select the ridge regularization parameter within the given ESN specification. The candidate with the lowest criterion value is then selected for forecasting. Time series cross-validation (rolling origin forecasting) could provide an alternative basis for choosing λ, but would substantially increase the computational cost given the number of series and ESN

configurations considered. Details on the information criteria are provided in Appendix A.2.

4 Empirical analysis

4.1 Accuracy metrics

Many error metrics are available in the forecasting literature to evaluate the accuracy of forecasting methods. Here, the out-of-sample forecasting performance is measured by the symmetric mean absolute percentage error (sMAPE), the mean absolute scaled error (MASE), and the overall weighted average (OWA). The sMAPE and MASE are defined as

$$sMAPE = \frac{2}{h} \sum_{t=T+1}^{T+h} \frac{|y_t - \hat{y}_t|}{|y_t| + |\hat{y}_t|} 100\%, \tag{6}$$

$$MASE = \frac{1}{h} \frac{\sum_{t=T+1}^{T+h} |y_t - \hat{y}_t|}{\frac{1}{T-m} \sum_{t=m+1}^{T} |y_t - y_{t-m}|}, \tag{7}$$

where y_t is the actual value and $\hat{y}_t$ is the estimated point forecast for a given period, h is the forecast horizon, T is the number of observations available in-sample, and m is the frequency of the time series, i.e., $m = 12$ for monthly and $m = 4$ for quarterly data (Hyndman and Athanasopoulos, 2021; Hyndman and Koehler, 2006; Makridakis et al., 2020).

The sMAPE calculates the absolute difference between the actual value and the forecast (the numerator in equation (6)). It divides by half the sum of the absolute value of the actuals and the absolute value of the forecasts (the denominator in equation (6)). This value is then summed for all points in time and divided by the forecast horizon.

In contrast to other percentage metrics, the sMAPE treats positive and negative forecast errors equally (symmetric). Percentage error metrics cannot be used when there are zero or close-to-zero values because there would be a division by zero, or the values would tend to infinity. Thus, the metric can be unstable. In the present application, however, this concern is substantially reduced because the original M4 dataset was specifically prepared to exclude negative, zero, and near-zero values. The sMAPE is still a helpful error metric, as it is easy to interpret.

The MASE is one of the most widely used forecast error metrics in contemporary forecasting literature. It overcomes the issues of the other metrics mentioned above and is more suitable from a research perspective. The idea of MASE is to scale the out-of-sample mean absolute error (MAE) of the forecasting method at hand (the numerator in equation (7)) with the in-sample MAE of some benchmark method (the denominator in equation (7)). In the original publication of Hyndman and Köhler (Hyndman and Koehler, 2006), the naive forecast method ($m = 1$) is proposed as the benchmark. However, benchmark methods such as the seasonal naive approach ($m = 12$ or $m = 4$) provide a more appropriate benchmark for seasonal time series (Makridakis et al., 2020).

The MASE has some appealing properties when comparing the accuracy of different methods and time series. For example, the MASE is scale-invariant, meaning that it is independent of the scale of the data. Hence, it allows for comparisons of forecasts across different scales (e.g., revenue in millions EUR vs. revenue in thousands EUR). There are no problems with zero or close-to-zero values because there is no division by the actual values, and the MASE is symmetric in that positive and negative forecast errors are treated equally. Finally, MASE is smaller than one if the examined forecast method performs better than the benchmark, and MASE is larger than one when the benchmark method performs better (Hyndman and Athanasopoulos, 2021; Hyndman and Koehler, 2006; Makridakis et al., 2020).

To provide a combined measure of forecast accuracy, the M4 Forecasting Competition uses the OWA, which averages the relative sMAPE and MASE of a forecasting method compared with the Naive2 benchmark (Makridakis et al., 2020). For model j, it is defined as

$$OWA_j = \frac{1}{2}\left(\frac{\overline{sMAPE}_j}{\overline{sMAPE}_{\text{NAIVE2}}} + \frac{\overline{MASE}_j}{\overline{MASE}_{\text{NAIVE2}}}\right), \tag{8}$$

where $\overline{sMAPE}_j$ and $\overline{MASE}_j$ denote the mean accuracy values of model j across all time series. By construction, Naive2 has an OWA of one. Values below one indicate that a method performs better overall than Naive2, while values above one indicate worse overall performance. The OWA therefore provides a single ranking measure that gives equal weight to the relative sMAPE and MASE results.

4.2 Hyperparameter sweep

To systematically explore the effect of hyperparameter configurations, we perform a comprehensive hyperparameter sweep based on an exhaustive grid search. The grid spans four core ESN parameters: the leakage rate α, the spectral radius ρ, the reservoir scaling factor τ and the information criterion. The leakage rate is varied in equidistant steps from 0.1 to 1.0, allowing the model to interpolate between slowly adapting reservoirs with long memory and rapidly updating reservoirs that prioritize short-term dynamics. The spectral radius is swept from 0.2 to 1.2 in increments of 0.1, covering both strictly contractive regimes as well as values slightly above unity, which are known to place the reservoir near the edge of stability and can empirically enhance predictive performance in certain settings. The reservoir size is explored via the scaling parameter $\tau \in \{0.2, 0.4, 0.6\}$, which allows the capacity of the model to adjust to the length of each time series while maintaining a fixed upper bound to avoid excessive complexity. Each configuration was further evaluated using information criteria including Akaike Information Criterion (AIC) (Akaike, 1974), corrected Akaike Information Criterion (AICc) (Sugiura, 1978; Hurvich and Tsai, 1989), Bayesian Information Criterion (BIC) (Schwarz, 1978), and Hannan-Quinn Criterion (HQC) (Hannan and Quinn, 1979) to account for model parsimony relative to goodness of fit. These criteria provide a complementary perspective by penalizing overly complex reservoir configurations, thereby supporting a more balanced model selection beyond pure predictive performance.

The resulting grid comprises 1,320 distinct hyperparameter configurations ($10 \times 11 \times 3 \times 4$), each of which is evaluated independently for every time series. This yields a total of 4,752,000 ESN fits across both frequencies (1,320 configurations across 2,400 monthly and 1,200 quarterly series), which remains computationally feasible due to the closed-form training of the linear readout. The chosen search space is deliberately compact yet expressive: it covers qualitatively distinct dynamical regimes while avoiding unnecessarily fine-grained sampling that would not lead to practically meaningful differences in model behavior. By evaluating each configuration independently on each time series, the procedure accommodates the heterogeneous temporal characteristics within the datasets and supports the selection of a frequency-specific hyperparameter configuration based on aggregate forecast accuracy.

To illustrate the qualitative effect of individual hyperparameters on the forecast dynamics, Figure 3 presents the results for the example series M21655 from the monthly dataset. The left panel shows forecasts obtained for different leakage rates α (ranging from 0.1 to 1.0), while the right panel visualizes forecasts for varying spectral radii ρ (ranging from 0.2 to 1.2). The actual holdout observations are plotted in black, and the forecasts are color-coded from blue to orange to reflect increasing hyperparameter values. As expected, low leakage values lead to smoother

and more inert forecast trajectories, whereas higher leakage rates produce more reactive but also more volatile forecasts. A similar pattern is observed for the spectral radius: small ρ values result in strongly damped dynamics, while larger values allow the reservoir to retain more temporal information that amplifies fluctuations. This visual example highlights how the hyperparameter configuration governs the responsiveness and memory characteristics of the ESN, motivating the need for a systematic search across a broad grid rather than relying on a single default setting.

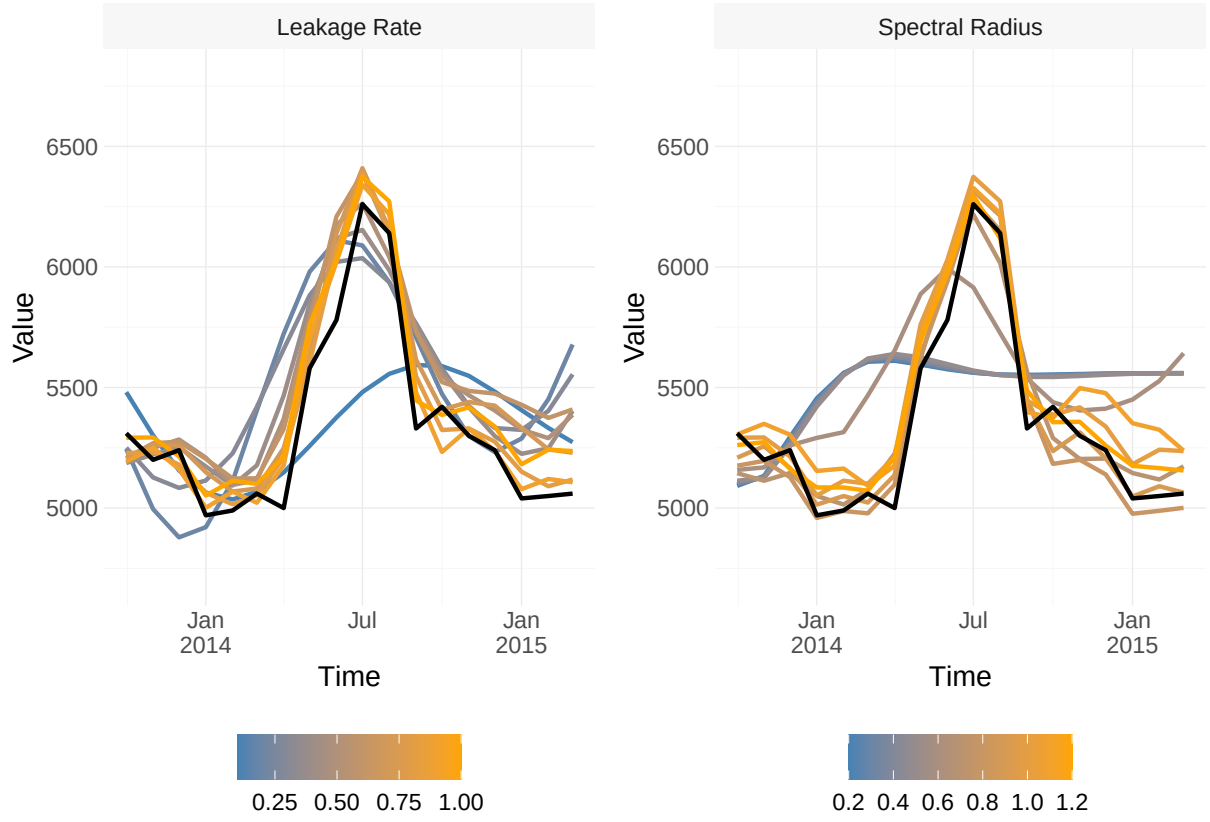

Figure 3: Forecasts for the example monthly series M21655 under varying hyperparameter settings. Left: forecasts for leakage rates $\alpha \in \{0.1, 0.2, ..., 0.9, 1.0\}$. Right: forecasts for spectral radii $\rho \in \{0.2, 0.3, ..., 1.1, 1.2\}$. Actual observations (holdout data) are shown in black. Colored forecast lines transition from blue (low values) to orange (high values), illustrating how increasing α and ρ alters forecast smoothness, reactivity, and temporal lag.

To complement the visualization-based analysis of individual hyperparameter effects, Tables 1 and 2 summarize the 30 best-performing ESN configurations for the monthly and quarterly datasets, respectively. By ranking configurations based on mean MASE and reporting both MASE and sMAPE alongside the corresponding hyperparameters and information criteria, these tables allow for a structured comparison of recurring patterns in optimal settings across frequencies.

Table 1 lists the 30 best-performing hyperparameter configurations for the monthly dataset, ranked by mean MASE. Across the top-ranked configurations, a clear pattern emerges: almost all high-performing models operate with a leakage rate of $\alpha = 0.9$ or $\alpha = 1.0$, indicating that more reactive reservoirs with limited smoothing tend to be favorable when sufficient data points are available, as is the case in monthly series of typical length. Similarly, the spectral radius con-

centrates around $\rho = 0.8$ to $\rho = 1.0$, placing the reservoir close to the edge of stability. This confirms the earlier observation from the visualization example that slightly more expressive and less damped dynamics improve forecasting accuracy. In terms of reservoir size scaling, $\tau = 0.4$ appears most frequently among the top-ranked configurations, with $\tau = 0.6$ occurring in some cases but never dominating. This suggests that moderately sized reservoirs are sufficient to capture relevant temporal structure, whereas further increasing the number of units does not consistently translate into better performance. Across all information criteria, AICc and HQC appear most frequently among the top configurations, but differences remain minor, showing that the selection of the information criteria is less important.

Table 2 shows the corresponding results for the quarterly dataset. Compared to the monthly results, the ranking structure is more homogeneous: almost all top configurations share $\alpha = 1.0$, confirming that when fewer observations per time series are available, the ESN benefits from updating the internal state as aggressively as possible rather than retaining long memory. The optimal spectral radius shifts towards smaller values around $\rho = 0.3$ to $\rho = 0.5$, indicating that more contractive dynamics are preferable in shorter quarterly time series to prevent overextension of the internal memory. In contrast to the monthly dataset, higher τ-values ($\tau = 0.6$) occur more systematically among the top-performing configurations. This suggests that, despite shorter time series, allocating a larger reservoir compensates for the reduced temporal resolution inherent to quarterly data. Interestingly, the dispersion across information criteria is again small, with AIC, AICc, HQC, and BIC all represented in the top ranks, reinforcing that the criterion for model selection is less critical. Overall, while both datasets favor high leakage values, the optimal spectral radius and reservoir scaling clearly depend on data frequency and series length, supporting the decision to perform a full grid search rather than fixing hyperparameters a priori.

To obtain a global view of hyperparameter sensitivity across all series, Figure 4 summarizes the median MASE values for each hyperparameter setting, computed over all models in the grid for both datasets. The panels are organized by hyperparameter dimension, covering leakage rate α, spectral radius ρ, reservoir scaling τ, and the information criterion. Within each panel, each point represents the median MASE for a specific hyperparameter value while marginalizing over all remaining settings, which yields smooth trajectories due to aggregation over a large number of time series and model configurations. Monthly and quarterly results are shown using a consistent color scheme (orange for monthly, blue for quarterly). The optimal value per panel, defined as the minimum median MASE, is indicated by a filled marker, whereas all remaining values are shown as hollow markers. Tables 11 and 12 in the Appendix report the corresponding numerical values.

For the monthly dataset, the lowest median MASE is achieved at $\alpha = 0.9$, $\rho = 0.7$, $\tau = 0.4$, and model selection based on BIC. For the quarterly dataset, the best performing values shift toward $\alpha = 1.0$ and $\rho = 0.3$, while $\tau = 0.4$ remains unchanged and BIC again yields the lowest median MASE. These patterns suggest frequency-specific differences in the preferred reservoir dynamics: quarterly time series tend to favor more contractive dynamics and more aggressive updating, whereas monthly data favor slightly more persistent reservoirs. At the same time, the differences between information criteria are small, suggesting that the choice among AIC, AICc, BIC, and HQC has limited practical impact in this application. Overall, the smoothness of the curves suggests stable empirical tendencies within the considered grid, rather than sharply separated optimal settings.

In summary, the results indicate broadly consistent and interpretable effects of ESN hyperparameters on forecasting performance, while also showing that the optimal configurations are not universal across frequencies. While fast state updates are preferred in both datasets, differences in spectral radius and reservoir capacity highlight the importance of adapting ESN dynamics to the temporal resolution of the data. These findings support the design choice of conducting a hy-

Rank	Model	Model Specification				MASE		sMAPE [%]	
		IC	α	ρ	τ	Mean	Median	Mean	Median
1	ESN-0650	AICc	1.0	0.9	0.4	**0.878**	0.712	17.592	12.406
2	ESN-1310	HQC	1.0	0.9	0.4	0.878	0.715	17.601	12.438
3	ESN-0320	AIC	1.0	0.9	0.4	0.881	0.715	17.701	12.509
4	ESN-0980	BIC	1.0	0.9	0.4	0.881	0.721	17.631	12.522
5	ESN-0317	AIC	1.0	0.8	0.4	0.883	0.713	17.599	12.600
6	ESN-0321	AIC	1.0	0.9	0.6	0.884	0.723	17.766	12.735
7	ESN-0647	AICc	1.0	0.8	0.4	0.884	**0.711**	17.593	12.600
8	ESN-0651	AICc	1.0	0.9	0.6	0.884	0.725	17.743	12.623
9	ESN-1277	HQC	0.9	0.9	0.4	0.884	0.713	17.594	12.571
10	ESN-0318	AIC	1.0	0.8	0.6	0.885	0.719	17.745	12.618
11	ESN-0617	AICc	0.9	0.9	0.4	0.885	0.713	17.594	12.587
12	ESN-0648	AICc	1.0	0.8	0.6	0.885	0.718	17.723	12.532
13	ESN-0947	BIC	0.9	0.9	0.4	0.885	0.713	17.610	12.560
14	ESN-1278	HQC	0.9	0.9	0.6	0.885	0.722	17.674	12.695
15	ESN-1308	HQC	1.0	0.8	0.6	0.885	0.716	17.685	12.556
16	ESN-1311	HQC	1.0	0.9	0.6	0.885	0.723	17.735	12.604
17	ESN-1313	HQC	1.0	1.0	0.4	0.885	0.716	17.890	12.314
18	ESN-0618	AICc	0.9	0.9	0.6	0.886	0.721	17.671	12.763
19	ESN-0653	AICc	1.0	1.0	0.4	0.886	0.719	17.903	**12.254**
20	ESN-0948	BIC	0.9	0.9	0.6	0.886	0.720	17.687	12.661
21	ESN-0983	BIC	1.0	1.0	0.4	0.886	0.720	17.898	12.337
22	ESN-1307	HQC	1.0	0.8	0.4	0.886	0.715	17.647	12.644
23	ESN-0287	AIC	0.9	0.9	0.4	0.887	0.713	17.629	12.692
24	ESN-0981	BIC	1.0	0.9	0.6	0.887	0.721	17.732	12.623
25	ESN-0285	AIC	0.9	0.8	0.6	0.888	0.718	17.644	12.593
26	ESN-0288	AIC	0.9	0.9	0.6	0.888	0.721	17.736	12.741
27	ESN-1275	HQC	0.9	0.8	0.6	0.889	0.717	17.672	12.509
28	ESN-1280	HQC	0.9	1.0	0.4	0.889	0.714	17.869	12.465
29	ESN-0614	AICc	0.9	0.8	0.4	0.890	0.720	**17.590**	12.791
30	ESN-0615	AICc	0.9	0.8	0.6	0.890	0.718	17.673	12.584

Table 1: **Out-of-sample accuracy** for the 30 best ESN configurations evaluated on **2,400 monthly time series** in the dataset *Parameter*. For every model, we show the information criterion, leakage rate α, spectral radius ρ, the reservoir size scaling τ, and the resulting mean/median values of MASE and sMAPE. The rows are sorted by the lowest mean MASE. The minimum value in each accuracy column is highlighted in bold, and the corresponding rows are shaded in light gray.

Rank	Model	Model Specification				MASE		sMAPE [%]	
		IC	α	ρ	τ	Mean	Median	Mean	Median
1	ESN-0306	AIC	1.0	0.4	0.6	**1.078**	0.875	10.243	5.491
2	ESN-0636	AICc	1.0	0.4	0.6	1.078	0.871	10.234	5.492
3	ESN-1296	HQC	1.0	0.4	0.6	1.079	0.871	10.235	5.495
4	ESN-0305	AIC	1.0	0.4	0.4	1.081	**0.861**	10.240	5.430
5	ESN-0309	AIC	1.0	0.5	0.6	1.081	0.868	10.267	5.491
6	ESN-0966	BIC	1.0	0.4	0.6	1.081	0.869	10.248	5.518
7	ESN-1295	HQC	1.0	0.4	0.4	1.081	0.863	10.228	5.441
8	ESN-1299	HQC	1.0	0.5	0.6	1.081	0.870	10.254	5.482
9	ESN-0303	AIC	1.0	0.3	0.6	1.082	0.869	10.244	5.457
10	ESN-0633	AICc	1.0	0.3	0.6	1.082	0.869	10.242	5.479
11	ESN-0635	AICc	1.0	0.4	0.4	1.082	0.861	10.244	5.450
12	ESN-0639	AICc	1.0	0.5	0.6	1.082	0.873	10.263	5.482
13	ESN-1298	HQC	1.0	0.5	0.4	1.082	0.872	**10.222**	5.447
14	ESN-1301	HQC	1.0	0.6	0.4	1.082	0.870	10.233	5.469
15	ESN-0308	AIC	1.0	0.5	0.4	1.083	0.864	10.243	5.440
16	ESN-0638	AICc	1.0	0.5	0.4	1.083	0.864	10.242	5.440
17	ESN-0968	BIC	1.0	0.5	0.4	1.083	0.867	10.231	5.463
18	ESN-0969	BIC	1.0	0.5	0.6	1.083	0.873	10.267	5.535
19	ESN-0972	BIC	1.0	0.6	0.6	1.083	0.875	10.271	5.503
20	ESN-0975	BIC	1.0	0.7	0.6	1.083	0.871	10.252	5.522
21	ESN-1293	HQC	1.0	0.3	0.6	1.083	0.869	10.248	5.454
22	ESN-1302	HQC	1.0	0.6	0.6	1.083	0.869	10.273	5.514
23	ESN-0311	AIC	1.0	0.6	0.4	1.084	0.863	10.251	5.495
24	ESN-0632	AICc	1.0	0.3	0.4	1.084	0.862	10.251	5.429
25	ESN-0642	AICc	1.0	0.6	0.6	1.084	0.865	10.264	5.519
26	ESN-0645	AICc	1.0	0.7	0.6	1.084	0.869	10.272	5.528
27	ESN-0971	BIC	1.0	0.6	0.4	1.084	0.875	10.240	5.490
28	ESN-0978	BIC	1.0	0.8	0.6	1.084	0.867	10.273	5.480
29	ESN-1304	HQC	1.0	0.7	0.4	1.084	0.865	10.265	5.612
30	ESN-0302	AIC	1.0	0.3	0.4	1.085	0.862	10.262	**5.415**

Table 2: **Out-of-sample accuracy** for the 30 best ESN configurations evaluated on **1,200 quarterly time series** in the dataset *Parameter*. For every model, we show the information criterion, leakage rate α, spectral radius ρ, the reservoir size scaling τ, and the resulting mean/median values of MASE and sMAPE. The rows are sorted by the lowest mean MASE. The minimum value in each accuracy column is highlighted in bold, and the corresponding rows are shaded in light gray.

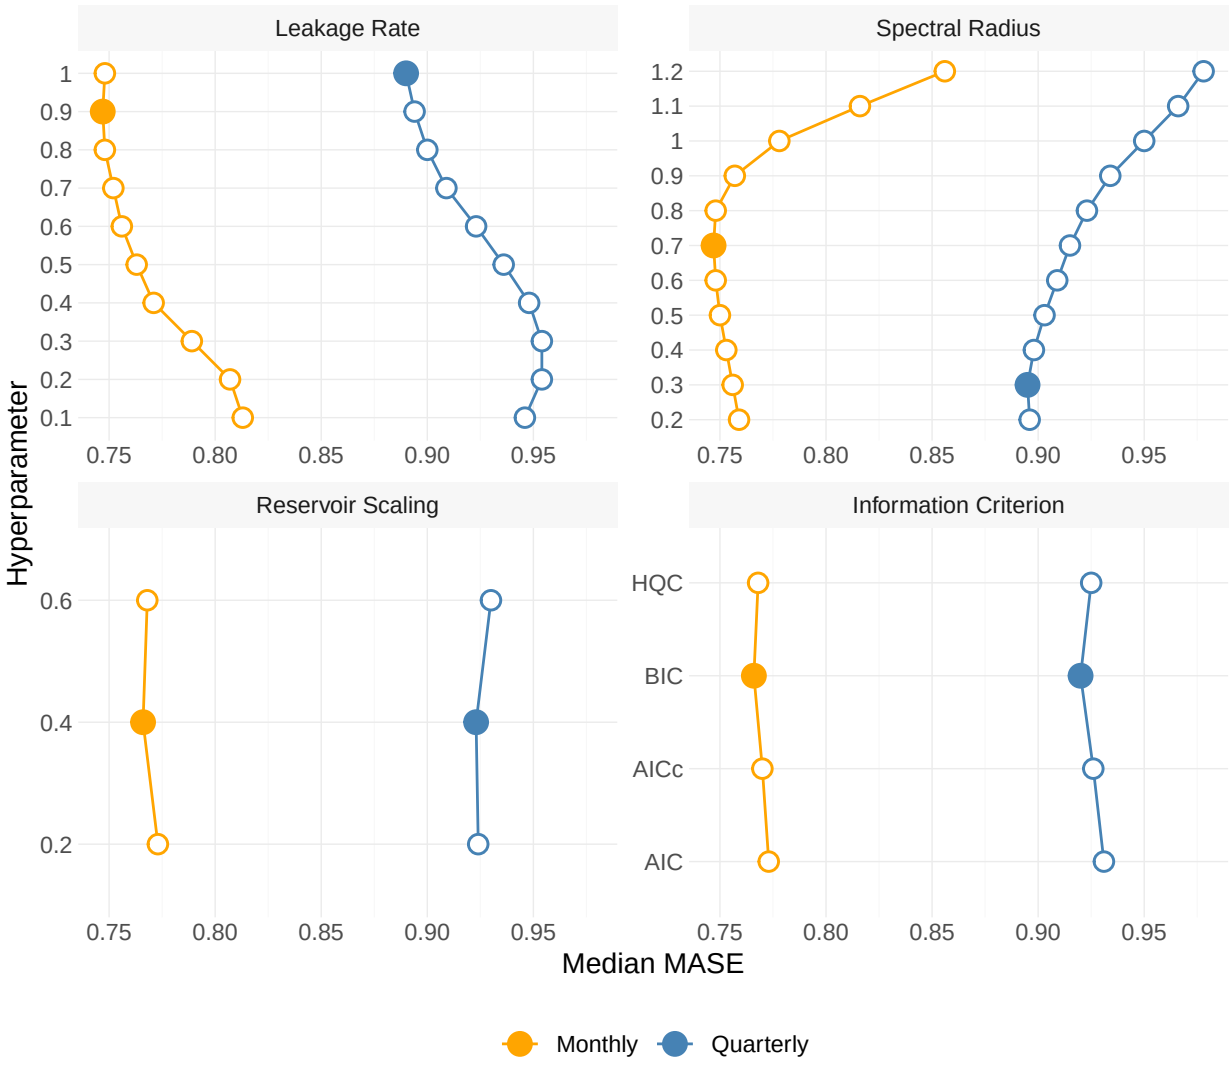

Figure 4: Median MASE across all time series as a function of individual hyperparameter settings, computed by marginalizing over all remaining grid configurations. Monthly results are shown in orange and quarterly results in blue. Each point denotes the median MASE for a specific hyperparameter value; the best-performing value per panel (minimum median MASE) is indicated by a filled marker, while all other values are shown as hollow markers.

perparameter sweep and provide empirical guidance for selecting ESN configurations in similar forecasting settings.

4.3 Random initialization

Although the ESN hyperparameters determine the overall reservoir dynamics, the model remains stochastic because the reservoir weights are randomly initialized (Lukoševičius and Jaeger, 2009; Lukoševičius, 2012). Different initializations generate different reservoir-state representations and may therefore lead to different fitted readout weights and recursive forecast paths, even when the data, preprocessing steps, and hyperparameters are held constant. Assessing this source of

variation is important because the results of a single ESN fit should not depend excessively on an arbitrary random seed.

To examine the sensitivity of the selected ESN specifications to random initialization, the best-performing monthly and quarterly configurations from the hyperparameter sweep are re-estimated using 30 different random seeds. For the monthly data, all models use AICc, $\alpha = 1.0$, $\rho = 0.9$, and $\tau = 0.4$. For the quarterly data, all models use AIC, $\alpha = 1.0$, $\rho = 0.4$, and $\tau = 0.6$. Thus, the models denoted by ESN-01 to ESN-30 differ only in their random initialization. Seed 42, which was used in the original hyperparameter sweep, is retained as the reference initialization. Each model is evaluated on the complete frequency-specific *Parameter* dataset using MASE and sMAPE.

Tables 3 and 4 summarize the resulting forecast accuracy and rank the initializations by mean MASE. For the monthly data, mean MASE ranges from 0.878 to 0.907, while median MASE ranges from 0.707 to 0.730. The corresponding mean sMAPE values range from 17.592% to 17.908%, and median sMAPE ranges from 12.406% to 13.074%. Seed 42 achieves the lowest mean MASE and the lowest mean and median sMAPE, whereas seed 229 produces the lowest median MASE. However, the differences between most initializations are small: excluding the lowest-ranked initialization, mean MASE lies between 0.878 and 0.894.

The comparatively high mean MASE of 0.907 for seed 840 is primarily caused by an isolated extreme error rather than a general deterioration across the monthly series. As reported in Appendix Table 13, this initialization produces a maximum MASE of 60.403 and a standard deviation of 1.387. For all other monthly initializations, the maximum MASE ranges from 7.110 to 8.483 and the standard deviation from 0.673 to 0.698. At the same time, seed 840 achieves a mean sMAPE of 17.641%, which is among the lowest values in Table 3. This difference illustrates that aggregate mean accuracy can occasionally be affected by a small number of initialization-specific extreme errors and reinforces the value of considering both mean and median results.

The quarterly results show even less sensitivity to random initialization. Mean MASE varies only between 1.078 and 1.089, while median MASE ranges from 0.860 to 0.884. Mean sMAPE lies between 10.227% and 10.324%, and median sMAPE between 5.459% and 5.578%. Seeds 42 and 622 achieve the lowest mean MASE after rounding, while seed 602 provides the lowest median MASE. Seed 623 achieves the lowest mean sMAPE, and seed 900 the lowest median sMAPE. Consequently, no single initialization dominates across all accuracy measures. The complete distributions in Appendix Table 14 are also highly similar: the maximum MASE varies only from 7.234 to 7.301 and its standard deviation from 0.824 to 0.838 across the 30 initializations.

Because the selected hyperparameter configurations were originally identified using seed 42 on the same *Parameter* dataset, the strong ranking of this reference seed should not be interpreted as evidence that seed 42 is intrinsically superior. Instead, the relevant finding is the narrow dispersion of forecast accuracy across the alternative seeds. In particular, the quarterly results and the monthly median results indicate that the aggregate performance of the selected ESN configurations is largely stable with respect to random initialization.

Figure 5 provides a complementary series-level comparison. For each time series, the MASE obtained from an initialization is expressed as a percentage deviation from the median MASE across all tested initializations for that series. This normalization controls for differences in forecasting difficulty between time series. The seed-specific median deviations remain concentrated around the zero-reference line for both frequencies, indicating that no initialization produces a consistently large improvement or deterioration across the complete dataset. Negative values identify seeds that perform slightly better than the series-specific median on average, whereas positive values indicate slightly worse performance. The limited spread around zero is consistent with the narrow ranges reported in Tables 3 and 4.

Rank	Model	Seed	MASE		sMAPE [%]	
			Mean	Median	Mean	Median
1	ESN-01	42	0.878	0.712	17.592	12.406
2	ESN-23	900	0.880	0.713	17.699	12.450
3	ESN-02	562	0.885	0.720	17.692	12.680
4	ESN-06	75	0.885	0.710	17.819	12.735
5	ESN-07	229	0.885	0.707	17.750	12.498
6	ESN-19	533	0.885	0.720	17.693	12.878
7	ESN-08	147	0.886	0.716	17.687	12.599
8	ESN-10	50	0.886	0.725	17.751	12.681
9	ESN-25	985	0.886	0.723	17.673	12.664
10	ESN-26	284	0.886	0.725	17.888	12.848
11	ESN-29	518	0.886	0.716	17.728	12.508
12	ESN-30	213	0.886	0.726	17.779	12.854
13	ESN-11	129	0.887	0.716	17.741	12.542
14	ESN-05	154	0.888	0.723	17.863	12.727
15	ESN-03	998	0.889	0.717	17.754	12.669
16	ESN-22	880	0.889	0.719	17.767	12.886
17	ESN-24	298	0.889	0.723	17.868	12.708
18	ESN-27	933	0.889	0.726	17.815	12.743
19	ESN-20	411	0.890	0.727	17.865	12.698
20	ESN-04	322	0.891	0.721	17.850	12.932
21	ESN-12	304	0.891	0.714	17.814	12.616
22	ESN-16	602	0.891	0.725	17.835	12.726
23	ESN-21	883	0.892	0.730	17.898	13.074
24	ESN-13	24	0.893	0.721	17.859	12.560
25	ESN-28	622	0.893	0.717	17.819	12.495
26	ESN-09	635	0.894	0.729	17.877	12.488
27	ESN-15	357	0.894	0.714	17.803	12.636
28	ESN-17	166	0.894	0.722	17.908	12.625
29	ESN-18	623	0.894	0.725	17.863	12.701
30	ESN-14	840	0.907	0.717	17.641	12.627

Table 3: **Out-of-sample accuracy** of the ESN across 30 random initializations evaluated on **2,400 monthly time series** in the dataset *Parameter*. For each initialization, the table reports the random seed and the resulting mean and median values of MASE and sMAPE. The rows are sorted by mean MASE from low to high.

While aggregate accuracy is comparatively stable, individual forecast paths can still differ across reservoir realizations. Figure 6 illustrates this effect for the monthly series M21655 and the quarterly series Q13921. All preprocessing choices and frequency-specific hyperparameters are fixed, and only the random initialization is varied. The resulting forecast paths are therefore not identical and may differ in their level, responsiveness, and development over the recursive forecast horizon. This variation arises because each reservoir generates a different nonlinear representation of the same input history. The figure demonstrates that random initialization can matter for an individual time series even when its effect largely averages out across a large collection

Rank	Model	Seed	MASE		sMAPE [%]	
			Mean	Median	Mean	Median
1	ESN-01	42	1.078	0.875	10.243	5.491
2	ESN-28	622	1.078	0.871	10.241	5.488
3	ESN-04	322	1.079	0.866	10.253	5.480
4	ESN-09	635	1.079	0.867	10.253	5.500
5	ESN-19	533	1.079	0.863	10.252	5.505
6	ESN-02	562	1.080	0.869	10.249	5.533
7	ESN-27	933	1.080	0.867	10.249	5.479
8	ESN-11	129	1.081	0.864	10.254	5.476
9	ESN-05	154	1.082	0.866	10.252	5.578
10	ESN-16	602	1.082	0.860	10.281	5.463
11	ESN-18	623	1.082	0.864	10.227	5.462
12	ESN-25	985	1.082	0.862	10.253	5.472
13	ESN-08	147	1.083	0.875	10.277	5.493
14	ESN-14	840	1.083	0.874	10.264	5.560
15	ESN-20	411	1.083	0.877	10.274	5.505
16	ESN-12	304	1.084	0.869	10.262	5.478
17	ESN-13	24	1.084	0.868	10.261	5.517
18	ESN-15	357	1.084	0.867	10.264	5.515
19	ESN-07	229	1.085	0.872	10.263	5.530
20	ESN-10	50	1.085	0.866	10.311	5.508
21	ESN-17	166	1.085	0.866	10.299	5.510
22	ESN-06	75	1.086	0.866	10.294	5.483
23	ESN-21	883	1.086	0.864	10.287	5.478
24	ESN-22	880	1.086	0.871	10.308	5.484
25	ESN-23	900	1.086	0.870	10.303	5.459
26	ESN-24	298	1.086	0.871	10.324	5.556
27	ESN-30	213	1.086	0.868	10.282	5.489
28	ESN-03	998	1.088	0.871	10.284	5.555
29	ESN-26	284	1.088	0.884	10.311	5.523
30	ESN-29	518	1.089	0.871	10.273	5.503

Table 4: **Out-of-sample accuracy** of the ESN across 30 random initializations evaluated on **1,200 quarterly time series** in the dataset *Parameter*. For each initialization, the table reports the random seed and the resulting mean and median values of MASE and sMAPE. The rows are sorted by mean MASE from low to high.

of series.

Overall, the analysis indicates that the selected ESN specifications are reasonably robust to random initialization at the aggregate level. The exact accuracy values and forecasts for individual series may change with the random seed, and isolated unfavorable reservoir realizations can produce large errors. Nevertheless, the small differences in mean and median accuracy across most seeds suggest that the main empirical conclusions are not driven by a broadly unstable initialization. Fixing the random seed remains necessary for reproducibility, while averaging forecasts across several independently initialized reservoirs could provide an additional way to reduce the

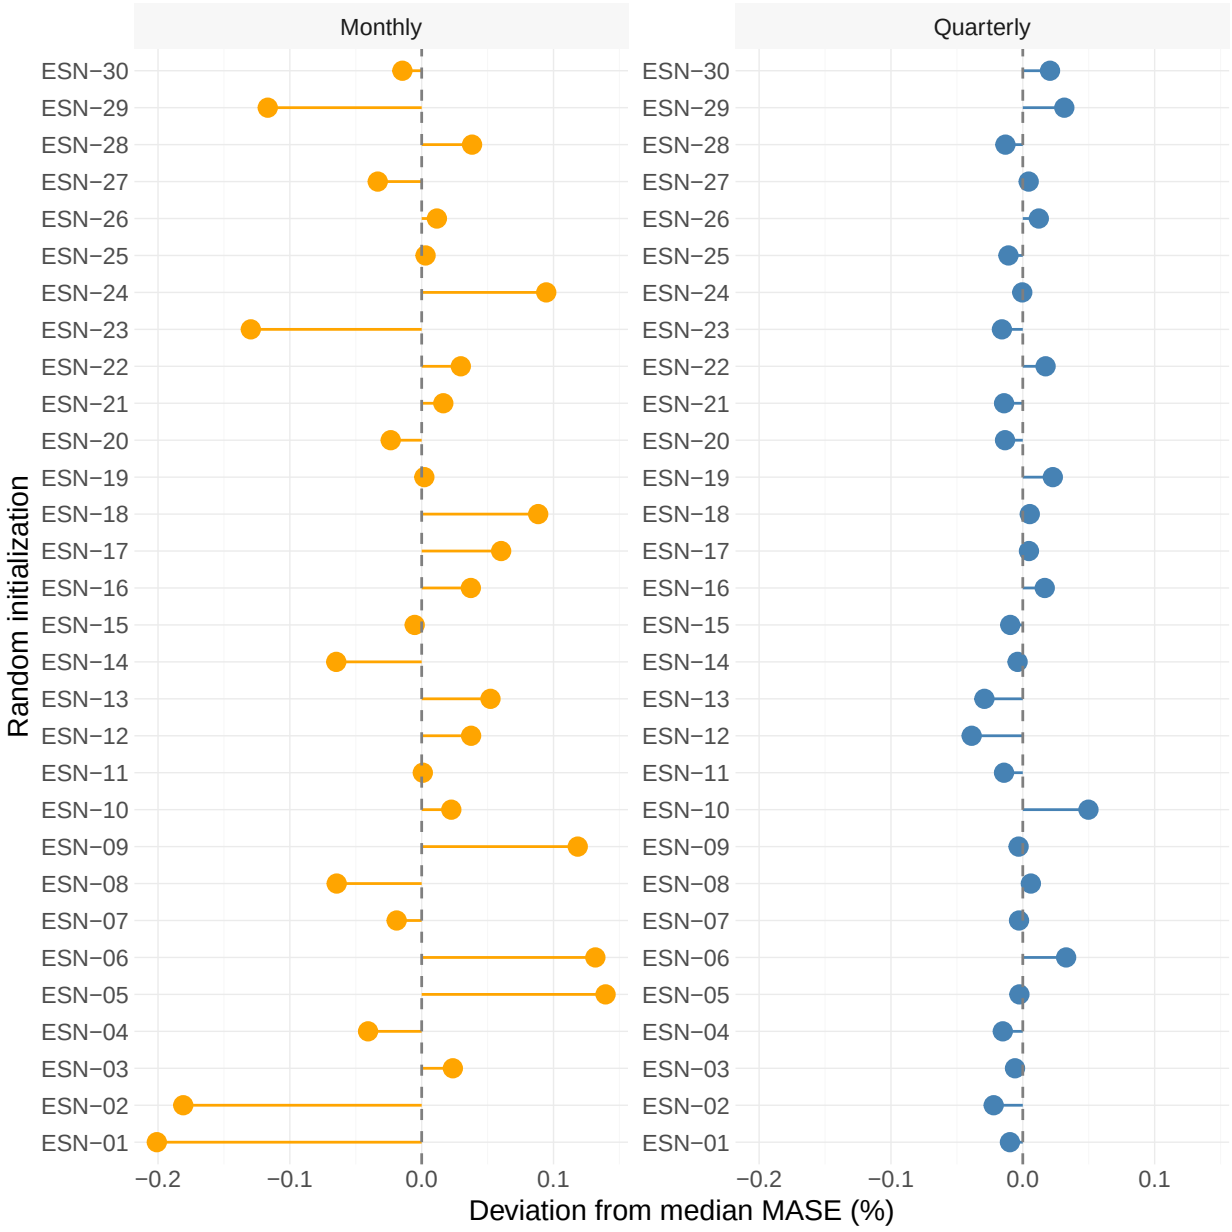

Figure 5: Median deviation from the series-specific median MASE across random initializations. For each time series, the MASE obtained from a given initialization is expressed as a percentage deviation from the median MASE across all tested initializations. Monthly results are shown in orange and quarterly results in blue. Each point represents the median deviation across all time series for a specific random initialization. Negative values indicate better-than-median forecast accuracy, whereas positive values indicate worse-than-median forecast accuracy.

remaining initialization-specific variation (i.e., ensemble forecasting).

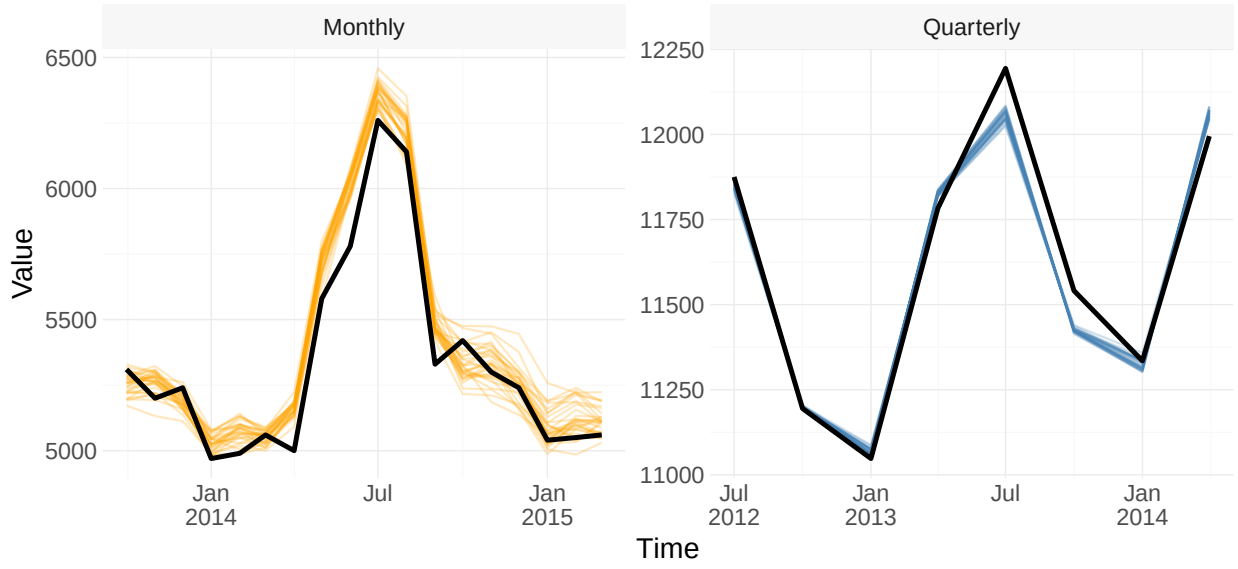

Figure 6: Forecasts under varying random initializations of the ESN. Left: forecasts for the example monthly series M21655 over an 18-period forecast horizon. Right: forecasts for the example quarterly series Q13921 over an 8-period forecast horizon. All frequency-specific hyperparameters are held constant, while only the random seed used to initialize the reservoir is varied. Actual observations (holdout data) are shown in black. Monthly forecasts are shown in orange and quarterly forecasts in blue, illustrating the sensitivity of the resulting forecasts to random initialization.

4.4 Benchmark analysis

In order to evaluate the performance of ESNs, forecast accuracy and computational runtime are compared against simple, statistical, and NN forecasting methods. The choice of benchmark methods is mainly driven by their relevance to automated forecasting and their use in forecasting research and practice. In addition to established simple and statistical approaches, the comparison includes the MLP and RNN benchmarks from the M4 Forecasting Competition and the competition winning model ES-RNN.

The reported ESN runtime refers to training and forecasting on the *Forecast* dataset after the main ESN hyperparameters α, ρ, and τ, as well as the information criterion, have been selected on the separate *Parameter* dataset. The ridge penalty λ, however, is still selected within each fitted ESN by random search using the chosen information criterion. The reported runtime therefore includes ESN training, the internal selection of λ, and recursive forecasting for the selected ESN specification, but not the preceding large-scale hyperparameter sweep on the *Parameter* dataset. Importantly, the MLP, RNN, and ES-RNN forecasts were taken from the published M4 benchmark and competition results and were not re-trained in the computational environment used for this study. Consequently, the runtime comparison is restricted to the ESN and the locally estimated simple and statistical benchmark methods and should not be interpreted as a complete computational comparison across all evaluated forecasting methods.

Table 5 summarizes the simple, statistical, and NN methods used in the forecast evaluation. The simple methods provide transparent reference points with different assumptions about the underlying time series. The seasonal naive method, for example, reproduces the most recent observation from the corresponding season and can perform well when seasonality is strong and other systematic changes are limited. Naive2 extends the naive approach by first removing sea-

sonality when it is statistically detected and subsequently restoring the estimated seasonal pattern to the forecasts. The drift method allows the most recent observation to increase or decrease according to the average historical change and may therefore be more appropriate for trending series (Hyndman and Athanasopoulos, 2021; Makridakis et al., 1998, 2020).

The statistical benchmarks include automatic ARIMA and ETS models, which are widely used in empirical forecasting studies (Hyndman et al., 2008; Hyndman and Khandakar, 2008; Makridakis et al., 2020), as well as the Theta method (Assimakopoulos and Nikolopoulos, 2000) and TBATS (Livera et al., 2011). In addition, the analysis includes the MLP and RNN benchmarks used in the M4 Forecasting Competition. Both NN benchmarks are fitted separately to each time series using lagged observations after detrending and seasonal adjustment, with forecasts generated recursively over the required horizon (Makridakis et al., 2020). Finally, ES-RNN is included as a more advanced hybrid benchmark. This method combines series-specific exponential smoothing components with an long short-term memory (LSTM) network trained across multiple time series and constituted the winning method of the M4 Forecasting Competition (Smyl, 2020).

The simple and statistical benchmark methods are estimated locally, whereas the MLP, RNN, and ES-RNN forecasts are taken from the corresponding published M4 benchmark and competition results (Makridakis et al., 2020; Smyl, 2020) and were not re-trained locally. Consequently, comparable computational runtimes are not available for these three methods. Details on the software packages, computational workflow, and computing environment are provided in Appendix A.5.

In the benchmark analysis, the forecast accuracy of the benchmark methods listed in Table 5 is compared with the best-performing ESN configurations identified in the hyperparameter sweep. Specifically, these are ESN-0650 for the monthly data (AICc, $\alpha = 1.0$, $\rho = 0.9$, $\tau = 0.4$) and ESN-0306 (AIC, $\alpha = 1.0$, $\rho = 0.4$, $\tau = 0.6$) for the quarterly data.

Tables 6 and 7 summarize the results for the monthly and quarterly datasets regarding forecast accuracy and computational runtime. Forecast accuracy is reported using OWA and the arithmetic mean and median of MASE and sMAPE across all time series of the respective frequency. Computational runtime is reported as total runtime for the entire monthly or quarterly dataset and average runtime per time series, measured in seconds. Runtime values are not available for MLP, RNN, and ES-RNN because their forecasts are taken from the published M4 benchmark and competition results. The models are ranked by OWA, with lower values indicating better overall forecast accuracy. The best value in each column is highlighted in bold, and the ESN is marked in light gray.

Table 6 reports the forecasting performance of 13 models on 2,400 monthly univariate time series. ES-RNN achieves the lowest OWA (0.858) and the best value for all four reported accuracy statistics, with a mean MASE of 0.870, a median MASE of 0.679, a mean sMAPE of 15.746%, and a median sMAPE of 11.579%. Among the statistical models, TBATS attains the lowest OWA (0.898), followed by Theta (0.902), ARIMA (0.914), and ETS (0.916). ESN-0650 ranks sixth overall with an OWA of 0.918. Nevertheless, its mean MASE of 0.898 is nearly identical to those of ARIMA (0.897) and TBATS (0.899), while its mean sMAPE of 17.463% is higher than the values obtained by the leading methods. Regarding computational efficiency, Naive2 and the other simple methods have the lowest runtimes, whereas TBATS is by far the slowest method. ESN-0650 requires approximately 0.36 seconds per time series and is faster than ARIMA and TBATS, but slower than ETS and the simple benchmarks.

The monthly results indicate that the ESN remains competitive with the established statistical models, particularly under MASE. Although it does not lead the OWA ranking in Table 6, the differences between ESN-0650, ARIMA, and TBATS in mean MASE are very small. Its lower OWA rank is mainly attributable to its comparatively weaker sMAPE performance, since OWA

Model	Description
NAIVE	Naive forecast (random walk model). All forecasts are set to the value of the last observation of the training data.
NAIVE2	Naive forecast applied to seasonally adjusted data when seasonality is detected. The forecasts are subsequently transformed back by restoring the estimated seasonal pattern.
DRIFT	Naive forecast plus drift (random walk with drift). A variation of the naive method, where the forecast is allowed to increase or decrease over time. The amount of change over time (drift) is set to be the average change in the training data.
SNAIVE	Seasonal naive forecast. Each forecast is set to be equal to the last observed value from the same season (e.g., for monthly data the same month of the previous year).
MEAN	Mean forecast. All forecasts are set to the arithmetic mean (average) of the training data.
ARIMA	Automatic ARIMA model using the Hyndman-Khandakar algorithm. Searches through the model space ARIMA$(p, d, q)(P, D, Q)_m$ to identify the best seasonal or nonseasonal ARIMA model with the lowest AICc value.
ETS	Automatic exponential smoothing state space model, based on the classification of exponential smoothing methods (Error, Trend, Season). For example, ETS(M, A, A) is Holt-Winters' method with multiplicative error, additive trend and additive seasonality. Model estimation is done via log-likelihood and model selection via AICc.
THETA	Theta method of Assimakopoulos and Nikolopoulos (2000). Theta is equivalent to simple exponential smoothing (ETS(A, N, N)) with drift applied to a nonseasonal or seasonally adjusted series.
TBATS	Exponential smoothing state space model with trigonometric seasonality, Box-Cox transformation, ARMA errors, trend and seasonal components. Originally intended for high-frequency time series with multiple and nested seasonal patterns.
MLP	Multilayer perceptron benchmark with a single feed-forward hidden layer. The model uses lagged, detrended, and seasonally adjusted observations as inputs and generates forecasts recursively.
RNN	Simple recurrent neural network benchmark with a recurrent hidden layer and linear output. The model uses lagged, detrended, and seasonally adjusted observations and generates forecasts recursively.
ES-RNN	Hybrid exponential smoothing and recurrent neural network model. Exponential smoothing represents series-specific level and seasonality, while an LSTM network learns common nonlinear patterns across time series.

Table 5: **Benchmark methods:** Overview and description of the simple, statistical, and neural network methods used for the comparison of forecast accuracy.

combines the relative mean MASE and mean sMAPE results. ES-RNN, however, outperforms the ESN and all other evaluated methods across the reported aggregate accuracy measures. The computational profile of the ESN remains favorable relative to ARIMA and especially TBATS. Overall,

the monthly results support the ESN as a competitive and computationally efficient alternative to established statistical methods, but not as the best-performing method once ES-RNN and the combined OWA measure are considered.

Model	OWA	MASE		sMAPE [%]		Runtime [sec]	
		Mean	Median	Mean	Median	Mean	Total
ES-RNN	**0.858**	**0.870**	**0.679**	**15.746**	**11.579**	–	–
TBATS	0.898	0.899	0.718	16.686	11.845	0.997	2393.176
THETA	0.902	0.911	0.724	16.638	11.666	0.036	86.347
ARIMA	0.914	0.897	0.733	17.334	11.845	0.411	986.246
ETS	0.916	0.905	0.713	17.257	11.696	0.279	668.774
ESN-0650	0.918	0.898	0.723	17.463	12.109	0.364	872.552
NAIVE2	1.000	0.999	0.803	18.635	13.097	**0.027**	**64.965**
NAIVE	1.076	1.110	0.859	19.386	14.387	0.029	70.492
DRIFT	1.080	1.101	0.835	19.716	13.263	0.029	70.788
SNAIVE	1.152	1.195	0.967	20.642	15.602	0.029	70.682
MLP	1.599	1.544	1.147	30.791	18.725	–	–
RNN	1.617	1.514	1.145	32.041	18.959	–	–
MEAN	2.462	2.893	1.992	37.797	33.115	0.028	68.006

Table 6: **Monthly dataset:** OWA, mean and median MASE and sMAPE, and computational runtime for the forecasting models evaluated on the 2,400 monthly time series in the dataset *Forecast*. The models are ranked according to OWA from low to high. The best-performing method (minimum value) per column is highlighted in bold, and the ESN is shaded in light gray. Runtime values are unavailable for ES-RNN, MLP, and RNN.

Table 7 presents forecasting accuracy results for 13 models evaluated on 1,200 quarterly time series. ES-RNN achieves the lowest OWA (0.883) and the best value for all four reported accuracy statistics, with a mean MASE of 1.097, a median MASE of 0.813, a mean sMAPE of 10.502%, and a median sMAPE of 5.344%. ETS ranks second according to OWA (0.900), followed by ESN-0306 (0.910), ARIMA (0.917), and TBATS (0.921). ESN-0306 achieves the second-lowest mean MASE (1.111) and the second-lowest median sMAPE (5.701%), while ETS provides a slightly lower mean sMAPE and median MASE. Regarding computational efficiency, the simple benchmarks are the fastest methods for which runtimes are available. TBATS is the slowest method, while ESN-0306 requires approximately 0.14 seconds per time series and has a runtime similar to ARIMA.

The quarterly results again demonstrate that the ESN yields competitive accuracy relative to the statistical benchmarks. It ranks third according to OWA in Table 7 and performs particularly well under mean MASE and median sMAPE. However, ES-RNN performs better across all reported aggregate accuracy measures, while ETS achieves a slightly lower OWA than the ESN. The results therefore show that the ESN is among the strongest methods considered for the quarterly data, although ES-RNN and ETS achieve better aggregate performance according to OWA. In terms of computational efficiency, the ESN is substantially faster than TBATS, although its runtime is slightly higher than those of ARIMA and ETS. Overall, the quarterly results indicate a favorable balance between forecast accuracy and computational cost, while also demonstrating the advantage of the more advanced globally trained ES-RNN model.

Figure 7 complements the descriptive accuracy results with pairwise statistical comparisons based on the MASE values for the individual time series. For each pair of forecasting methods, a

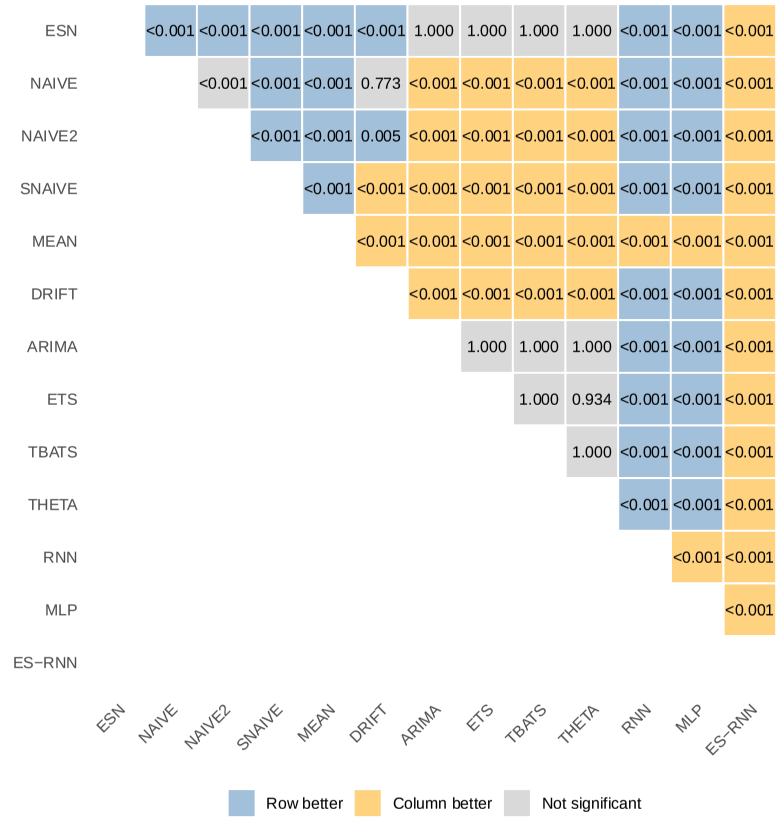

(a) Monthly dataset

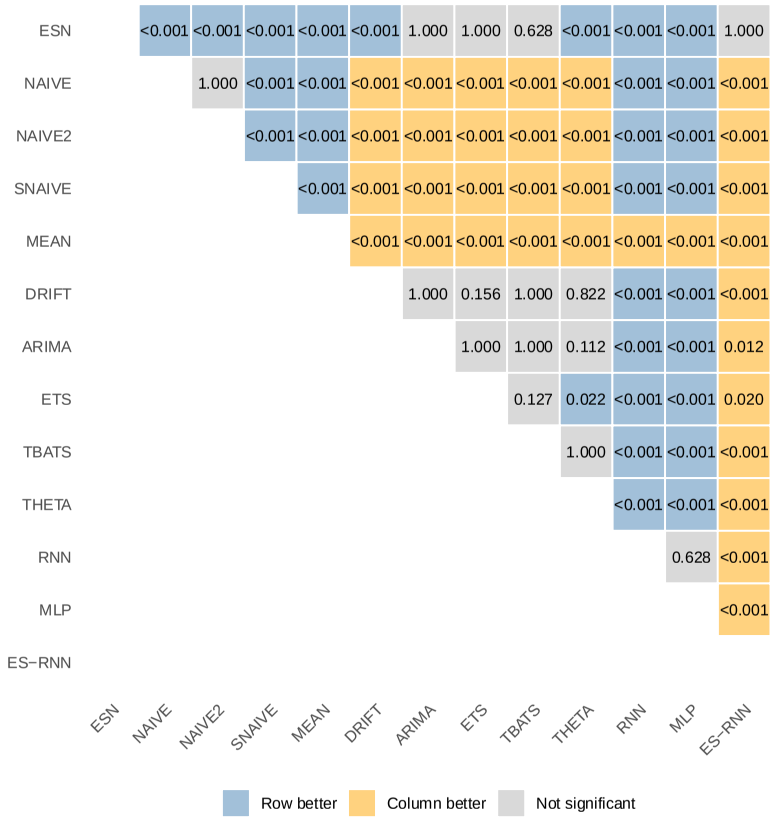

(b) Quarterly dataset

Figure 7: Pairwise comparisons of forecasting models based on MASE. The cells report Holm-adjusted p-values from paired Wilcoxon signed-rank tests across the individual time series. Blue cells indicate that the row model has significantly lower MASE than the column model, orange cells indicate that the column model has significantly lower MASE, and gray cells indicate that no statistically significant difference is detected.

Model	OWA	MASE		sMAPE [%]		Runtime [sec]	
		Mean	Median	Mean	Median	Mean	Total
ES-RNN	**0.883**	**1.097**	**0.813**	**10.502**	**5.344**	–	–
ETS	0.900	1.121	0.868	10.689	5.805	0.076	91.462
ESN-0306	0.910	1.111	0.875	11.023	5.701	0.136	163.243
ARIMA	0.917	1.139	0.872	10.917	5.875	0.123	147.634
TBATS	0.921	1.160	0.879	10.832	6.028	0.505	605.496
THETA	0.930	1.171	0.922	10.932	6.186	0.032	38.478
DRIFT	0.972	1.192	0.908	11.706	6.005	0.028	33.692
NAIVE2	1.000	1.286	1.027	11.519	6.697	0.027	32.683
NAIVE	1.040	1.347	1.082	11.888	7.145	0.032	38.175
SNAIVE	1.164	1.522	1.281	13.185	8.194	0.028	34.031
MLP	1.512	1.935	1.454	17.505	9.866	–	–
RNN	1.539	1.995	1.485	17.584	10.048	–	–
MEAN	2.974	4.240	3.521	30.522	25.888	**0.027**	**32.320**

Table 7: **Quarterly dataset:** OWA, mean and median MASE and sMAPE, and computational runtime for the forecasting models evaluated on the 1,200 quarterly time series in the dataset *Forecast*. The models are ranked according to OWA from low to high. The best-performing method (minimum value) per column is highlighted in bold, and the ESN is shaded in light gray. Runtime values are unavailable for ES-RNN, MLP, and RNN.

paired Wilcoxon signed-rank test (Wilcoxon, 1945) is used to examine whether the distribution of the series-level MASE differences is centered around zero. Because multiple pairwise comparisons are conducted, the resulting p-values are adjusted using the Holm procedure (Holm, 1979). The tests therefore assess whether the observed differences between methods are systematic across the time series rather than being driven only by differences in aggregate mean accuracy.

For the monthly data, the ESN achieves significantly lower MASE than the naive method, Naive2, the seasonal naive method, the mean and drift forecasts, as well as the MLP and RNN benchmarks. By contrast, no statistically significant difference is detected between the ESN and ARIMA, ETS, TBATS, or the Theta method. These results support the conclusion drawn from the descriptive statistics that the ESN performs at a level comparable to the strongest statistical benchmarks for monthly forecasting. The ES-RNN is the only model that achieves significantly lower monthly MASE than the ESN. However, the corresponding median difference is comparatively small, indicating that statistical significance does not necessarily imply a large practical difference in forecast accuracy.

For the quarterly data, the ESN significantly outperforms the naive, Naive2, seasonal naive, mean, drift, Theta, MLP, and RNN methods. No statistically significant differences are found relative to ARIMA, ETS, TBATS, or ES-RNN. Thus, although the ESN achieves the second-lowest mean MASE in Table 7, the pairwise tests do not establish that it is systematically more accurate than these closely competing models across the individual quarterly series. Taken together with Figure 7, the results indicate that the ESN clearly improves upon the simpler benchmarks and the standalone neural-network models, while its performance is generally statistically indistinguishable from that of the strongest statistical methods and ES-RNN. The exact median differences and Holm-adjusted p-values for comparisons involving the ESN are reported in Appendix Tables 19 and 20.

Overall, the results in Tables 6 and 7 show that the proposed autoregressive ESN is empirically competitive with established statistical forecasting methods in the considered setting. Its performance is particularly strong under MASE, whereas its ranking is less favorable under sMAPE and the combined OWA measure. ES-RNN achieves the best aggregate accuracy for both frequencies, while Figure 7 shows that the ESN is generally statistically indistinguishable from the strongest statistical benchmarks. Thus, the results support the ESN as a competitive and computationally efficient alternative to standard statistical methods, but not as a uniformly superior approach across all models, metrics, and frequencies. Additional detail is provided in Appendix Tables 15, 16, 17, and 18, which report the minimum, first quartile, mean, median, third quartile, maximum, and standard deviation of the forecast errors.

Before summarizing the main findings, several aspects of the empirical design should be kept in mind. The analysis is based on filtered monthly and quarterly M4 series with at most 20 years of history, so the results refer to this forecasting setting rather than to the full M4 benchmark. Moreover, the ESN specification is intentionally simple and uses only the first autoregressive lag as input; richer input structures, including explicit seasonal lags, may lead to different results. The comparison includes the MLP, RNN, and ES-RNN models from the M4 Forecasting Competition, but it does not provide a systematic evaluation of newer neural forecasting architectures or more advanced reservoir computing variants. Finally, although pairwise statistical comparisons are reported for MASE, the results remain conditional on the reservoir initialization scheme and the set of forecasting methods considered in this study.

5 Summary and concluding remarks

This paper investigated the use of ESNs for univariate time series forecasting in a large-scale setting. Building on the M4 Forecasting Competition dataset (Makridakis et al., 2020), we focused on monthly and quarterly series with at most 20 years of historical data to reflect realistic constraints found in many organizational environments, where long histories are often unavailable or only partially relevant. To ensure that the empirical analysis remained both rigorous and practically meaningful, the dataset was split into two randomly drawn subsets: a *Parameter* dataset for systematic ESN hyperparameter exploration and a disjoint *Forecast* dataset for out-of-sample evaluation against benchmark methods. An exploratory data analysis showed that both subsets preserve several important characteristics of the filtered monthly and quarterly M4 series, while also reflecting the imposed restriction on the time series length.

Methodologically, the study proposed a first-order autoregressive ESN framework tailored to univariate forecasting. The approach combines standard preprocessing steps like stationarity testing via the KPSS test, differencing when required, and scaling to a symmetric interval with a leaky integrator ESN whose reservoir size is scaled with the time series length and capped to avoid excessive complexity. The internal reservoir dynamics are governed by three key hyperparameters: leakage rate, spectral radius, and reservoir scaling. The linear readout is trained by ridge regression, where the regularization parameter is selected via random search guided by information criteria. This design provides a flexible yet conceptually simple forecasting pipeline that can be applied in a fully automated fashion across large collections of time series.

A central contribution of the paper is the comprehensive hyperparameter sweep, which systematically evaluated the impact of leakage rate, spectral radius, reservoir size, and information criterion across thousands of time series and more than four million ESN fits. The results show that ESN hyperparameters exert broadly consistent and interpretable effects on forecasting performance. For monthly data, high leakage rates ($\alpha \approx 0.9 - 1.0$), spectral radii close to the edge of

stability ($\rho \approx 0.8 - 1.0$), and moderate reservoir scaling ($\tau \approx 0.4$) yielded the best overall performance. For quarterly data, optimal configurations shifted towards even higher leakage ($\alpha = 1.0$) and more contractive dynamics (spectral radius around $\rho = 0.3 - 0.5$), while reservoir scaling tended to be somewhat larger. These patterns reflect differences in temporal resolution and effective memory requirements between monthly and quarterly time series. The smoothness of the hyperparameter response curves suggests that the identified settings reflect stable empirical tendencies within the considered grid rather than sharply separated optimal configurations.

The empirical benchmark analysis on the *Forecast* dataset further demonstrated that the proposed ESN can compete with well-established statistical models in terms of forecast accuracy while offering favorable computational properties. For the monthly series, the best ESN configuration achieved a mean MASE that was nearly identical to those of ARIMA and TBATS, although ES-RNN provided the best aggregate accuracy according to OWA, MASE, and sMAPE. For the quarterly series, the ESN achieved the second-lowest mean MASE and the second-lowest median sMAPE and ranked third according to OWA, behind ES-RNN and ETS. ES-RNN again provided the best aggregate accuracy, whereas the standalone MLP and RNN benchmarks performed less favorably. Pairwise Wilcoxon signed-rank tests showed that the ESN clearly outperformed the simpler benchmarks and the standalone NN models, while its MASE was generally statistically indistinguishable from that of the strongest statistical methods. Across both frequencies, the selected ESN specification could be trained and applied efficiently on the *Forecast* dataset. These reported runtimes represent deployment after the main ESN configuration has been selected and exclude the preceding hyperparameter sweep comprising 4,752,000 ESN fits on the *Parameter* dataset, for which a reliable total runtime was not recorded. In addition, the MLP, RNN, and ES-RNN models were not re-trained locally, and comparable runtimes are therefore not available for these methods. Consequently, the reported runtime results should not be interpreted as evidence that the complete ESN model-selection pipeline is computationally more efficient than all evaluated alternatives. Taken together, the results indicate that simple ESNs can offer a useful trade-off between forecast accuracy and deployment cost once the main hyperparameters have been selected.

At the same time, several limitations of the present study highlight opportunities for future research. First, the analysis is restricted to univariate, autoregressive ESNs without exogenous inputs. Many real-world forecasting problems involve explanatory variables such as prices, promotions, or macroeconomic indicators; extending the ESN framework to multivariate or input-driven architectures would therefore be a natural next step. Second, we focused on point forecasts only and evaluated performance using sMAPE, MASE, and OWA. In many applications, decision makers require full predictive distributions or interval forecasts; integrating probabilistic output layers or ensemble techniques into the ESN framework could address this need (Xu and Han, 2016; Yao and Wang, 2019). Third, the study considered a single class of reservoir architectures with $\tanh(\cdot)$ activation, fixed sparsity, and basic random initialization. Alternative reservoir designs, such as modular reservoirs (Babinec and Pospíšchal, 2012), as well as deep learning approaches (i.e., multiple reservoirs) (Gallicchio and Micheli, 2017; Gallicchio et al., 2018), may further enhance performance or stability and merit systematic investigation.

Moreover, while the present work employed an exhaustive grid search for core ESN hyperparameters and random search for the ridge penalty, more advanced optimization strategies could be explored. Bayesian optimization (Maat et al., 2018) or Particle Swarm Optimization (Chouikhi et al., 2017) might reduce computational effort while still identifying high-quality configurations, particularly when extending the approach to richer model classes or larger datasets. The benchmark comparison includes selected NN and hybrid methods from the M4 Forecasting Competition, but it does not provide a systematic evaluation of newer neural forecasting architectures. Finally, the empirical analysis focused on monthly and quarterly frequencies; applying the proposed

ESN framework to other M4 frequencies (e.g., daily or hourly) or to newer benchmark datasets would provide additional evidence on its generality and limitations, especially in settings with complex multiple seasonality or high-frequency noise.

In summary, the findings suggest that simple ESNs can be a useful addition to the forecasting toolbox when many relatively short monthly or quarterly series must be processed automatically. The ESN is competitive with the strongest statistical benchmarks, particularly under MASE, but ES-RNN achieves the best aggregate accuracy for both frequencies. The evidence therefore supports empirical competitiveness rather than uniform superiority. The conclusions are conditional on the filtered M4 subsets, the first-order autoregressive ESN design, the fixed preprocessing pipeline, and the selected hyperparameter grid.

Appendix

A.1 Data

Application Field	Parameter				Forecast			
	Monthly		Quarterly		Monthly		Quarterly	
	n	%	n	%	n	%	n	%
Demographic	159	6.6	88	7.3	127	5.3	73	6.1
Finance	693	28.9	210	17.5	696	29.0	214	17.8
Industry	532	22.2	254	21.2	561	23.4	271	22.6
Macro	413	17.2	251	20.9	396	16.5	223	18.6
Micro	590	24.6	330	27.5	602	25.1	345	28.8
Other	13	0.5	67	5.6	18	0.8	74	6.2
Total	2,400	100.0	1,200	100.0	2,400	100.0	1,200	100.0

Table 8: Absolute (n) and relative (%) frequencies by application field for the datasets *Parameter* and *Forecast* per frequency.

Measuring strength of trend and strength of seasonality

To measure the strength of trend and strength of seasonality, the time series is decomposed via STL (Seasonal and Trend decomposition using LOESS) into the additive model $y_t = T_t + S_t + R_t$, where T_t is the trend-cycle component, S_t is the seasonal component and R_t is the remainder component (Cleveland et al., 1990). For time series with a strong trend, the seasonally adjusted series should have much more variation than the remainder component, and therefore $\text{Var}(R_t)/\text{Var}(T_t + R_t)$ should be relatively small. The two variances should be approximately the same for time series with little or no trend. The strength of the trend can be defined as

$$F_T = \max\left(0, 1 - \frac{\text{Var}(R_t)}{\text{Var}(T_t + R_t)}\right), \tag{9}$$

which measures the trend's strength between 0 and 1. The minimal possible value of F_T is set equal to zero as the variance of the remainder can be larger than the variance of the seasonally adjusted series.

The strength of seasonality can be defined as

$$F_S = \max\left(0, 1 - \frac{\text{Var}(R_t)}{\text{Var}(S_t + R_t)}\right), \tag{10}$$

similar to the trend strength, but for the detrended time series instead of the seasonally adjusted data. A value for F_S close to zero indicates almost no seasonality, while a value close to one exhibits a strong seasonal pattern within the time series because $\text{Var}(R_t)$ will be much smaller than $\text{Var}(S_t + R_t)$. These measures are helpful when analyzing a large number of time series to get an overview of the time series characteristics (Hyndman and Athanasopoulos, 2021; O'Hara-Wild et al., 2026b).

Metric	Dataset	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
Length	Total	42.000	82.000	216.300	202.000	306.000	2794.000	137.406
	Parameter	54.000	69.000	127.965	102.000	184.000	240.000	62.435
	Forecast	46.000	69.000	129.592	104.000	186.000	240.000	62.391
Season	Total	0.000	0.192	0.415	0.326	0.634	0.999	0.277
	Parameter	0.000	0.174	0.339	0.268	0.469	0.994	0.234
	Forecast	0.000	0.170	0.341	0.263	0.478	0.994	0.242
Trend	Total	0.015	0.894	0.889	0.971	0.993	1.000	0.188
	Parameter	0.037	0.805	0.848	0.953	0.983	1.000	0.214
	Forecast	0.015	0.813	0.847	0.950	0.983	1.000	0.213

Table 9: Descriptive statistics of the **monthly dataset**, summarizing the number of observations, seasonal strength and trend strength per time series for the complete M4 dataset ($n = 48,000$) and the randomly sampled datasets *Parameter* and *Forecast* (each $n = 2,400$).

Metric	Dataset	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
Length	Total	16.000	62.000	92.254	88.000	115.000	866.000	51.130
	Parameter	17.000	39.000	55.522	58.000	70.000	80.000	15.524
	Forecast	20.000	39.000	56.038	61.500	70.000	80.000	15.590
Season	Total	0.000	0.148	0.335	0.248	0.454	0.999	0.259
	Parameter	0.000	0.138	0.312	0.222	0.394	0.995	0.250
	Forecast	0.000	0.131	0.304	0.221	0.391	0.995	0.245
Trend	Total	0.149	0.959	0.954	0.990	0.998	1.000	0.092
	Parameter	0.361	0.946	0.946	0.986	0.996	1.000	0.097
	Forecast	0.289	0.943	0.941	0.983	0.996	1.000	0.108

Table 10: Descriptive statistics of the **quarterly dataset**, summarizing the number of observations, seasonal strength and trend strength per time series for the complete M4 dataset ($n = 24,000$) and the randomly sampled datasets *Parameter* and *Forecast* (each $n = 1,200$).

A.2 Hyperparameter sweep

Calculation of information criteria

For each candidate ridge penalty λ, the fitted readout produces the residual sum of squares

$$RSS_\lambda = \left\| \mathbf{y} - \mathbf{X}\widehat{\mathbf{W}}_\lambda^{\text{out}} \right\|_2^2. \tag{11}$$

Let n denote the number of observations used to estimate the readout after differencing, lag construction, and removal of the washout period. The effective degrees-of-freedom df_λ are calculated from the ridge hat matrix as defined in Equation 5. The intercept is not penalized.

Assuming Gaussian readout residuals, the maximized log-likelihood used in the implementation is

$$\ell_\lambda = -\frac{n}{2}\left[\log(2\pi) + 1 + \log\left(\frac{RSS_\lambda}{n}\right)\right]. \tag{12}$$

Based on this log-likelihood and the effective degrees-of-freedom, the information criteria are calculated as

$$AIC_\lambda = -2\ell_\lambda + 2df_\lambda. \tag{13}$$

$$AICc_\lambda = AIC_\lambda + \frac{2df_\lambda(df_\lambda + 1)}{n - df_\lambda - 1}. \tag{14}$$

$$BIC_\lambda = -2\ell_\lambda + df_\lambda \log(n). \tag{15}$$

$$HQC_\lambda = -2\ell_\lambda + 2df_\lambda \log(\log(n)). \tag{16}$$

For each fixed ESN configuration and reservoir realization, these criteria are evaluated for all randomly drawn candidate values of λ. The candidate producing the lowest value of the selected information criterion is retained for forecasting.

A.3 Random initialization

Par.	Value	MASE							sMAPE [%]						
		Min.	Q_1	Mean	Median	Q_3	Max.	Std.	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
α	0.1	0.010	0.494	1.047	0.813	1.317	129.508	0.984	0.074	4.880	20.326	14.090	28.692	197.694	21.576
	0.2	0.014	0.490	1.034	0.807	1.308	117.064	1.014	0.070	4.761	20.375	14.035	28.418	200.000	22.137
	0.3	0.015	0.482	1.001	0.789	1.281	10.102	0.804	0.069	4.648	19.864	13.628	27.725	199.616	21.513
	0.4	0.012	0.471	0.973	0.771	1.251	14.790	0.777	0.071	4.553	19.040	13.323	26.820	199.206	20.057
	0.5	0.012	0.466	0.957	0.763	1.230	10.521	0.759	0.073	4.488	18.676	13.143	26.524	200.000	19.427
	0.6	0.007	0.462	0.944	0.756	1.211	10.011	0.746	0.073	4.434	18.452	13.053	26.274	198.338	19.109
	0.7	0.005	0.459	0.938	0.752	1.201	10.209	0.742	0.071	4.413	18.309	13.020	26.174	200.000	18.852
	0.8	0.008	0.457	0.935	0.748	1.196	17.255	0.745	0.070	4.398	18.240	12.918	26.048	200.000	18.782
	0.9	0.003	0.455	0.935	0.747	1.194	154.659	0.796	0.074	4.392	18.242	12.862	26.066	200.000	18.886
	1.0	0.005	0.454	0.937	0.748	1.200	36.473	0.757	0.073	4.384	18.312	12.847	26.057	200.000	19.020
ρ	0.2	0.005	0.460	0.969	0.759	1.239	10.586	0.808	0.104	4.468	18.334	13.065	26.200	189.054	18.632
	0.3	0.007	0.459	0.961	0.756	1.233	10.586	0.795	0.071	4.444	18.250	13.050	26.081	189.552	18.548
	0.4	0.007	0.457	0.954	0.753	1.225	10.621	0.782	0.069	4.428	18.198	13.032	26.028	189.901	18.514
	0.5	0.014	0.456	0.947	0.750	1.217	10.802	0.767	0.071	4.424	18.154	13.025	25.956	190.010	18.480
	0.6	0.003	0.455	0.940	0.748	1.206	11.095	0.754	0.070	4.417	18.117	13.007	25.986	189.932	18.470
	0.7	0.005	0.455	0.934	0.747	1.198	11.215	0.743	0.070	4.407	18.123	13.007	26.024	191.672	18.520
	0.8	0.015	0.456	0.932	0.748	1.194	11.157	0.736	0.071	4.422	18.247	13.056	26.167	193.542	18.730
	0.9	0.010	0.464	0.939	0.757	1.201	10.906	0.735	0.071	4.449	18.618	13.126	26.552	194.810	19.375
	1.0	0.010	0.479	0.967	0.778	1.232	117.064	0.867	0.072	4.562	19.545	13.387	27.456	198.901	21.019
	1.1	0.012	0.503	1.028	0.816	1.308	85.601	0.869	0.072	4.839	21.002	14.004	29.056	200.000	23.185
	1.2	0.011	0.525	1.099	0.856	1.386	154.659	1.078	0.070	5.036	22.231	14.644	30.612	200.000	24.736
τ	0.2	0.007	0.469	0.974	0.773	1.251	10.588	0.784	0.070	4.513	19.145	13.237	27.125	200.000	20.201
	0.4	0.005	0.468	0.966	0.766	1.233	129.508	0.873	0.069	4.497	18.890	13.220	26.709	200.000	19.861
	0.6	0.003	0.468	0.970	0.768	1.238	154.659	0.797	0.070	4.582	18.917	13.319	26.663	200.000	19.911
IC	AIC	0.005	0.471	0.977	0.773	1.247	154.659	0.840	0.069	4.574	19.116	13.328	26.940	200.000	20.212
	AICc	0.005	0.470	0.972	0.770	1.243	129.508	0.816	0.070	4.548	19.030	13.287	26.867	200.000	20.067
	BIC	0.003	0.466	0.964	0.766	1.234	129.508	0.808	0.071	4.484	18.840	13.191	26.702	200.000	19.743
	HQC	0.003	0.467	0.968	0.768	1.239	129.508	0.812	0.071	4.514	18.948	13.239	26.788	200.000	19.941

Table 11: Descriptive statistics, summarizing the distribution of out-of-sample forecast accuracy per hyperparameter for $n = 2,400$ **monthly time series**. The table lists the minimum, first quartile (Q_1), mean, median, third quartile (Q_3), maximum, and standard deviation for MASE and sMAPE (%). The rows are grouped by hyperparameter and ordered by value.

Par.	Value	MASE							sMAPE [%]						
		Min.	Q_1	Mean	Median	Q_3	Max.	Std.	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
α	0.1	0.028	0.534	1.185	0.946	1.556	7.153	0.914	0.079	2.664	11.006	5.854	12.916	188.934	14.906
	0.2	0.037	0.535	1.187	0.954	1.578	7.309	0.908	0.113	2.690	11.179	5.902	13.035	200.000	15.663
	0.3	0.025	0.534	1.186	0.954	1.574	7.392	0.903	0.074	2.702	11.142	5.913	13.051	198.300	15.609
	0.4	0.026	0.529	1.177	0.948	1.560	7.876	0.894	0.073	2.683	11.070	5.897	12.847	200.000	15.591
	0.5	0.024	0.526	1.165	0.936	1.545	9.053	0.887	0.067	2.651	10.953	5.855	12.670	195.642	15.432
	0.6	0.024	0.520	1.152	0.923	1.525	7.390	0.877	0.067	2.606	10.834	5.779	12.517	195.763	15.293
	0.7	0.024	0.515	1.140	0.909	1.507	8.060	0.868	0.067	2.549	10.729	5.740	12.322	191.572	15.166
	0.8	0.024	0.509	1.127	0.900	1.484	8.334	0.859	0.067	2.519	10.627	5.696	12.177	200.000	15.083
	0.9	0.015	0.504	1.118	0.894	1.466	10.447	0.854	0.058	2.480	10.579	5.654	12.116	190.702	15.109
	1.0	0.012	0.502	1.112	0.890	1.457	9.013	0.852	0.046	2.450	10.525	5.613	12.085	198.089	15.027
ρ	0.2	0.026	0.499	1.124	0.896	1.485	7.423	0.865	0.074	2.480	10.492	5.678	12.081	187.019	14.648
	0.3	0.028	0.502	1.125	0.895	1.484	7.416	0.863	0.077	2.495	10.497	5.679	12.089	187.877	14.608
	0.4	0.025	0.504	1.127	0.898	1.488	7.405	0.863	0.073	2.510	10.517	5.686	12.145	189.066	14.607
	0.5	0.024	0.507	1.130	0.903	1.491	7.390	0.865	0.067	2.520	10.555	5.697	12.220	188.848	14.644
	0.6	0.024	0.510	1.135	0.909	1.495	7.378	0.868	0.067	2.530	10.611	5.723	12.271	188.074	14.720
	0.7	0.025	0.514	1.141	0.915	1.506	7.367	0.871	0.077	2.560	10.686	5.745	12.382	187.595	14.841
	0.8	0.014	0.520	1.149	0.923	1.518	7.425	0.874	0.054	2.609	10.789	5.757	12.498	187.124	15.032
	0.9	0.012	0.529	1.161	0.934	1.533	9.017	0.880	0.046	2.639	10.933	5.823	12.658	192.977	15.331
	1.0	0.016	0.537	1.177	0.950	1.554	9.013	0.888	0.064	2.680	11.146	5.909	12.866	198.870	15.747
	1.1	0.026	0.550	1.204	0.966	1.589	10.447	0.914	0.073	2.755	11.495	6.022	13.278	200.000	16.502
	1.2	0.026	0.557	1.228	0.978	1.620	9.771	0.943	0.074	2.794	11.788	6.100	13.509	198.803	17.212
τ	0.2	0.024	0.520	1.153	0.924	1.518	9.017	0.884	0.067	2.599	10.768	5.825	12.474	188.014	14.957
	0.4	0.012	0.519	1.150	0.923	1.519	7.890	0.876	0.046	2.580	10.837	5.749	12.567	200.000	15.221
	0.6	0.016	0.524	1.161	0.930	1.537	10.447	0.887	0.064	2.600	10.988	5.801	12.641	200.000	15.687
IC	AIC	0.012	0.524	1.163	0.931	1.537	10.447	0.890	0.046	2.617	10.956	5.823	12.673	200.000	15.475
	AICc	0.012	0.522	1.156	0.926	1.529	9.053	0.883	0.046	2.599	10.882	5.797	12.582	200.000	15.331
	BIC	0.012	0.518	1.146	0.920	1.511	8.286	0.874	0.046	2.562	10.766	5.756	12.457	200.000	15.076
	HQC	0.012	0.521	1.154	0.925	1.524	9.053	0.882	0.046	2.592	10.853	5.781	12.548	200.000	15.281

Table 12: Descriptive statistics, summarizing the distribution of out-of-sample forecast accuracy per hyperparameter for $n = 1,200$ **quarterly time series**. The table lists the minimum, first quartile (Q_1), mean, median, third quartile (Q_3), maximum, and standard deviation for MASE and sMAPE (%). The rows are grouped by hyperparameter and ordered by value.

Seed	MASE							sMAPE [%]						
	Min.	Q_1	Mean	Median	Q_3	Max.	Std.	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
24	0.011	0.440	0.893	0.721	1.145	7.138	0.697	0.154	4.143	17.859	12.560	25.855	187.578	18.605
42	0.023	0.425	0.878	0.712	1.137	7.167	0.680	0.078	4.097	17.592	12.406	25.324	187.273	18.287
50	0.013	0.432	0.886	0.725	1.141	7.166	0.684	0.119	4.223	17.751	12.681	25.428	186.516	18.503
75	0.022	0.439	0.885	0.710	1.148	7.192	0.673	0.095	4.216	17.819	12.735	25.621	187.880	18.569
129	0.022	0.436	0.887	0.716	1.138	7.764	0.690	0.082	4.231	17.741	12.542	26.047	187.984	18.472
147	0.016	0.435	0.886	0.716	1.133	7.110	0.688	0.119	4.234	17.687	12.599	25.795	187.864	18.232
154	0.030	0.440	0.888	0.723	1.141	7.176	0.677	0.083	4.232	17.863	12.727	25.820	188.170	18.489
166	0.016	0.435	0.894	0.722	1.146	7.119	0.692	0.091	4.179	17.908	12.625	25.950	187.371	18.809
213	0.034	0.436	0.886	0.726	1.133	7.192	0.680	0.115	4.173	17.779	12.854	25.919	180.621	18.261
229	0.025	0.433	0.885	0.707	1.131	7.149	0.693	0.080	4.195	17.750	12.498	25.503	182.495	18.462
284	0.014	0.443	0.886	0.725	1.145	7.163	0.679	0.123	4.120	17.888	12.848	25.852	189.189	18.537
298	0.017	0.442	0.889	0.723	1.138	7.142	0.676	0.088	4.135	17.868	12.708	26.047	191.852	18.392
304	0.015	0.441	0.891	0.714	1.137	7.179	0.686	0.111	4.265	17.814	12.616	25.904	189.496	18.536
322	0.023	0.441	0.891	0.721	1.130	7.172	0.689	0.086	4.145	17.850	12.932	25.592	187.978	18.521
357	0.023	0.438	0.894	0.714	1.153	7.572	0.698	0.079	4.184	17.803	12.636	25.930	188.524	18.465
411	0.022	0.437	0.890	0.727	1.141	8.483	0.694	0.083	4.178	17.865	12.698	25.897	187.261	18.758
518	0.018	0.434	0.886	0.716	1.132	7.171	0.684	0.099	4.219	17.728	12.508	25.713	188.676	18.477
533	0.022	0.432	0.885	0.720	1.135	7.148	0.684	0.170	4.103	17.693	12.878	25.785	189.458	18.351
562	0.031	0.433	0.885	0.720	1.144	7.170	0.680	0.085	4.223	17.692	12.680	25.619	187.370	18.412
602	0.026	0.439	0.891	0.725	1.134	7.181	0.686	0.097	4.170	17.835	12.726	25.907	185.723	18.566
622	0.014	0.443	0.893	0.717	1.145	7.125	0.694	0.160	4.137	17.819	12.495	25.874	187.931	18.554
623	0.021	0.439	0.894	0.725	1.145	7.207	0.690	0.132	4.173	17.863	12.701	26.142	187.672	18.430
635	0.021	0.440	0.894	0.729	1.153	7.908	0.694	0.154	4.183	17.877	12.488	26.006	187.202	18.591
840	0.021	0.442	0.907	0.717	1.133	60.403	1.387	0.132	4.211	17.641	12.627	25.551	187.959	18.341
880	0.027	0.436	0.889	0.719	1.134	8.099	0.694	0.142	4.231	17.767	12.886	25.825	188.482	18.439
883	0.011	0.439	0.892	0.730	1.147	7.187	0.679	0.100	4.142	17.898	13.074	25.815	182.699	18.445
900	0.035	0.435	0.880	0.713	1.129	7.193	0.679	0.097	4.142	17.699	12.450	25.387	183.248	18.482
933	0.022	0.439	0.889	0.726	1.144	7.156	0.680	0.127	4.127	17.815	12.743	25.755	188.500	18.539
985	0.016	0.438	0.886	0.723	1.131	7.195	0.678	0.189	4.181	17.673	12.664	25.555	188.326	18.257
998	0.025	0.436	0.889	0.717	1.151	7.175	0.693	0.084	4.232	17.754	12.669	25.738	187.699	18.416

Table 13: Descriptive statistics summarizing the distribution of out-of-sample forecast accuracy across $n = 2,400$ **monthly time series** for each random initialization of the ESN. The table reports the minimum, first quartile (Q_1), mean, median, third quartile (Q_3), maximum, and standard deviation of MASE and sMAPE (%). The rows are ordered by the random seed.

Seed	MASE							sMAPE [%]						
	Min.	Q_1	Mean	Median	Q_3	Max.	Std.	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
24	0.029	0.484	1.084	0.868	1.423	7.261	0.834	0.116	2.376	10.261	5.517	11.924	186.042	14.674
42	0.030	0.488	1.078	0.875	1.417	7.285	0.828	0.155	2.357	10.243	5.491	11.796	185.617	14.688
50	0.032	0.495	1.085	0.866	1.417	7.249	0.833	0.157	2.367	10.311	5.508	11.830	186.203	14.739
75	0.032	0.493	1.086	0.866	1.427	7.301	0.830	0.092	2.375	10.294	5.483	11.843	185.597	14.687
129	0.030	0.492	1.081	0.864	1.423	7.260	0.837	0.145	2.360	10.254	5.476	11.793	186.114	14.736
147	0.031	0.488	1.083	0.875	1.427	7.274	0.824	0.143	2.390	10.277	5.493	11.751	185.585	14.684
154	0.031	0.486	1.082	0.866	1.425	7.271	0.838	0.162	2.342	10.252	5.578	11.969	185.948	14.669
166	0.032	0.494	1.085	0.866	1.428	7.288	0.831	0.160	2.372	10.299	5.510	11.853	184.373	14.698
213	0.032	0.495	1.086	0.868	1.422	7.299	0.828	0.124	2.426	10.282	5.489	11.798	185.334	14.649
229	0.031	0.488	1.085	0.872	1.424	7.263	0.831	0.115	2.404	10.263	5.530	11.958	185.324	14.641
284	0.031	0.492	1.088	0.884	1.430	7.244	0.833	0.151	2.421	10.311	5.523	11.890	185.924	14.686
298	0.031	0.488	1.086	0.871	1.429	7.276	0.828	0.157	2.381	10.324	5.556	11.962	185.392	14.707
304	0.032	0.495	1.084	0.869	1.424	7.267	0.836	0.140	2.366	10.262	5.478	11.872	185.856	14.667
322	0.028	0.488	1.079	0.866	1.431	7.291	0.828	0.078	2.380	10.253	5.480	11.742	185.909	14.691
357	0.031	0.490	1.084	0.867	1.421	7.263	0.835	0.125	2.365	10.264	5.515	11.720	185.168	14.705
411	0.031	0.484	1.083	0.877	1.425	7.260	0.829	0.155	2.361	10.274	5.505	11.786	186.554	14.703
518	0.030	0.492	1.089	0.871	1.428	7.250	0.837	0.119	2.428	10.273	5.503	11.856	186.591	14.680
533	0.031	0.486	1.079	0.863	1.416	7.280	0.831	0.118	2.364	10.252	5.505	11.774	185.583	14.673
562	0.032	0.492	1.080	0.869	1.425	7.262	0.829	0.127	2.375	10.249	5.533	11.692	185.487	14.657
602	0.032	0.490	1.082	0.860	1.437	7.286	0.829	0.118	2.366	10.281	5.463	11.842	186.318	14.731
622	0.029	0.486	1.078	0.871	1.406	7.251	0.830	0.086	2.372	10.241	5.488	11.794	185.401	14.691
623	0.032	0.490	1.082	0.864	1.422	7.234	0.831	0.165	2.384	10.227	5.462	11.676	186.011	14.652
635	0.032	0.491	1.079	0.867	1.408	7.283	0.829	0.150	2.303	10.253	5.500	11.596	186.178	14.726
840	0.031	0.480	1.083	0.874	1.419	7.288	0.833	0.133	2.355	10.264	5.560	12.004	185.581	14.673
880	0.032	0.491	1.086	0.871	1.432	7.271	0.829	0.116	2.397	10.308	5.484	11.733	186.141	14.726
883	0.031	0.488	1.086	0.864	1.424	7.291	0.834	0.139	2.453	10.287	5.478	11.878	185.533	14.684
900	0.031	0.491	1.086	0.870	1.430	7.262	0.833	0.150	2.453	10.303	5.459	11.945	185.924	14.705
933	0.032	0.486	1.080	0.867	1.420	7.268	0.825	0.157	2.340	10.249	5.479	11.715	186.079	14.676
985	0.030	0.492	1.082	0.862	1.428	7.281	0.828	0.084	2.384	10.253	5.472	11.733	184.770	14.644
998	0.031	0.491	1.088	0.871	1.428	7.254	0.833	0.157	2.369	10.284	5.555	11.858	185.502	14.666

Table 14: Descriptive statistics summarizing the distribution of out-of-sample forecast accuracy across $n = 1,200$ **quarterly time series** for each random initialization of the ESN. The table reports the minimum, first quartile (Q_1), mean, median, third quartile (Q_3), maximum, and standard deviation of MASE and sMAPE (%). The rows are ordered by the random seed.

A.4 Benchmark analysis

Rank	Model	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
1	ES-RNN	0.041	0.427	0.870	0.679	1.040	18.354	0.932
2	ARIMA	0.014	0.451	0.897	0.733	1.093	14.453	0.794
3	ESN-0650	0.037	0.438	0.898	0.723	1.117	13.837	0.794
4	TBATS	0.008	0.449	0.899	0.718	1.112	15.235	0.809
5	ETS	0.000	0.458	0.905	0.713	1.129	12.496	0.775
6	THETA	0.050	0.437	0.911	0.724	1.131	15.275	0.823
7	NAIVE2	0.024	0.500	0.999	0.803	1.254	15.548	0.839
8	DRIFT	0.036	0.470	1.101	0.835	1.396	17.578	1.061
9	NAIVE	0.000	0.527	1.110	0.859	1.384	16.099	1.029
10	SNAIVE	0.000	0.624	1.195	0.967	1.484	20.416	1.024
11	RNN	0.060	0.741	1.514	1.145	1.908	19.448	1.249
12	MLP	0.081	0.726	1.544	1.147	1.954	18.923	1.338
13	MEAN	0.123	1.115	2.893	1.992	3.565	24.049	2.656

Table 15: **Monthly MASE:** Descriptive statistics summarizing the distribution of MASE per forecasting model for the 2,400 monthly time series in the dataset *Forecast*. The models are ranked according to mean MASE from low to high, and the ESN is shaded in light gray.

Rank	Model	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
1	ES-RNN	0.159	4.000	15.746	11.579	22.668	121.376	15.336
2	THETA	0.179	4.294	16.638	11.666	24.675	170.829	16.136
3	TBATS	0.075	4.110	16.686	11.845	24.335	140.503	16.228
4	ETS	0.000	4.158	17.257	11.696	25.118	155.392	17.360
5	ARIMA	0.077	4.103	17.334	11.845	24.976	129.936	17.383
6	ESN-0650	0.257	4.215	17.463	12.109	25.025	177.725	17.706
7	NAIVE2	0.116	4.596	18.635	13.097	27.012	138.387	18.152
8	NAIVE	0.000	4.908	19.386	14.387	28.547	138.387	18.117
9	DRIFT	0.292	4.981	19.716	13.263	27.999	194.421	20.647
10	SNAIVE	0.000	6.471	20.642	15.602	29.903	137.976	18.114
11	MLP	0.271	8.038	30.791	18.725	38.992	200.000	34.501
12	RNN	0.363	7.956	32.041	18.959	40.357	200.000	37.554
13	MEAN	0.531	16.727	37.797	33.115	55.014	174.643	26.168

Table 16: **Monthly sMAPE:** Descriptive statistics summarizing the distribution of sMAPE per forecasting model for the 2,400 monthly time series in the dataset *Forecast*. The models are ranked according to mean sMAPE from low to high, and the ESN is shaded in light gray.

Rank	Model	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
1	ES-RNN	0.037	0.493	1.097	0.813	1.426	9.634	0.962
2	ESN-0306	0.048	0.504	1.111	0.875	1.428	9.049	0.930
3	ETS	0.069	0.495	1.121	0.868	1.446	9.439	0.944
4	ARIMA	0.055	0.508	1.139	0.872	1.436	8.946	0.969
5	TBATS	0.055	0.524	1.160	0.879	1.449	9.522	1.014
6	THETA	0.043	0.550	1.171	0.922	1.496	9.048	0.948
7	DRIFT	0.058	0.529	1.192	0.908	1.523	9.885	1.041
8	NAIVE2	0.074	0.614	1.286	1.027	1.655	9.134	0.989
9	NAIVE	0.070	0.609	1.347	1.082	1.750	9.573	1.079
10	SNAIVE	0.097	0.746	1.522	1.281	1.977	9.850	1.108
11	MLP	0.059	0.825	1.935	1.454	2.593	42.974	1.917
12	RNN	0.073	0.823	1.995	1.485	2.584	99.922	3.224
13	MEAN	0.053	1.436	4.240	3.521	6.267	18.407	3.195

Table 17: **Quarterly MASE:** Descriptive statistics summarizing the distribution of MASE per forecasting model for the 1,200 quarterly time series in the dataset *Forecast*. The models are ranked according to mean MASE from low to high, and the ESN is shaded in light gray.

Rank	Model	Min.	Q_1	Mean	Median	Q_3	Max.	Std.
1	ES-RNN	0.098	2.670	10.502	5.344	11.535	154.973	14.438
2	ETS	0.156	2.685	10.689	5.805	12.186	166.740	14.342
3	TBATS	0.128	2.743	10.832	6.028	12.629	167.089	14.175
4	ARIMA	0.147	2.758	10.917	5.875	12.839	129.042	13.953
5	THETA	0.117	3.076	10.932	6.186	12.434	164.638	14.528
6	ESN-0306	0.177	2.590	11.023	5.701	12.272	168.859	15.533
7	NAIVE2	0.332	3.641	11.519	6.697	12.771	167.510	14.879
8	DRIFT	0.213	2.645	11.706	6.005	13.125	167.840	16.796
9	NAIVE	0.332	3.652	11.888	7.145	13.402	167.842	15.071
10	SNAIVE	0.359	4.599	13.185	8.194	15.556	153.666	14.840
11	MLP	0.133	4.499	17.505	9.866	21.747	200.000	21.734
12	RNN	0.182	4.559	17.584	10.048	21.649	189.089	21.700
13	MEAN	0.654	14.176	30.522	25.888	41.340	140.623	22.379

Table 18: **Quarterly sMAPE:** Descriptive statistics summarizing the distribution of sMAPE per forecasting model for the 1,200 quarterly time series in the dataset *Forecast*. The models are ranked according to mean sMAPE from low to high, and the ESN is shaded in light gray.

A.5 Computational details

The empirical analysis is conducted in the open-source programming language R (R Core Team, 2025). The complete computational workflow is organized in the R project `forecast_engine`, which contains the scripts for data preparation, construction of the *Parameter* and *Forecast* datasets, the ESN hyperparameter sweep, the random initialization analysis, benchmark estimation, forecast evaluation, and the generation of tables and figures.

The ESN models are estimated and forecast using the R package `echos` (Häußer, 2026a). The

Model	Median difference	Adjusted p-value	Decision
NAIVE	-0.106	< 0.001	ESN better
NAIVE2	-0.059	< 0.001	ESN better
SNAIVE	-0.175	< 0.001	ESN better
MEAN	-1.223	< 0.001	ESN better
DRIFT	-0.043	< 0.001	ESN better
ARIMA	0.006	1.000	Not significant
ETS	0.006	1.000	Not significant
TBATS	0.008	1.000	Not significant
THETA	0.012	1.000	Not significant
RNN	-0.361	< 0.001	ESN better
MLP	-0.338	< 0.001	ESN better
ES-RNN	0.032	< 0.001	Benchmark better

Table 19: Pairwise comparisons of MASE between the ESN and the benchmark models for the monthly series. Median differences are defined as $\text{MASE}_{\text{ESN}} - \text{MASE}_{\text{Benchmark}}$, such that negative values favor the ESN. Adjusted p-values are obtained from paired Wilcoxon signed-rank tests with Holm correction across all model comparisons. Results classified as "ESN better" or "Benchmark better" indicate significantly lower MASE for the corresponding model at the 5% level.

Model	Median difference	Adjusted p-value	Decision
NAIVE	-0.145	< 0.001	ESN better
NAIVE2	-0.113	< 0.001	ESN better
SNAIVE	-0.281	< 0.001	ESN better
MEAN	-2.506	< 0.001	ESN better
DRIFT	-0.011	< 0.001	ESN better
ARIMA	0.005	1.000	Not significant
ETS	-0.006	1.000	Not significant
TBATS	-0.005	0.628	Not significant
THETA	-0.033	< 0.001	ESN better
RNN	-0.444	< 0.001	ESN better
MLP	-0.419	< 0.001	ESN better
ES-RNN	0.004	1.000	Not significant

Table 20: Pairwise comparisons of MASE between the ESN and the benchmark models for the quarterly series. Median differences are defined as $\text{MASE}_{\text{ESN}} - \text{MASE}_{\text{Benchmark}}$, such that negative values favor the ESN. Adjusted p-values are obtained from paired Wilcoxon signed-rank tests with Holm correction across all model comparisons. Results classified as "ESN better" indicate significantly lower MASE for the ESN at the 5% level.

package is available on CRAN, and its development version is maintained in the GitHub repository. Supporting functionality for time series processing, model estimation, and forecast evaluation is provided by the R package `tscv` (Häußer, 2026b).

The simple and statistical benchmark methods are primarily estimated using the R package `fable` (O'Hara-Wild et al., 2026a). Since TBATS is not implemented in `fable`, this benchmark is estimated separately using the R package `forecast` (Hyndman et al., 2026). For the remaining sta-

tistical methods, the underlying forecasting algorithms are largely comparable across both packages. For example, `fable::ARIMA()` and `forecast::auto.arima()` both implement the Hyndman–Khandakar procedure for automatic ARIMA modeling (Hyndman and Khandakar, 2008). The MLP, RNN, and ES-RNN forecasts are taken from the corresponding M4 benchmark and competition forecasts (Makridakis et al., 2020; Smyl, 2020).

The empirical analysis was conducted using R version 4.5.2 on Windows 11 x64 (build 26200). All locally estimated models and runtime measurements were executed on a system equipped with an Intel Core i7-10510U CPU at 1.80 GHz and 16 GB RAM. The reported runtimes in Tables 6 and 7 refer to model estimation and forecast generation on the *Forecast* dataset. For the ESN, the runtime includes training, the internal random search used to select the ridge penalty λ, and recursive forecasting with the selected frequency-specific ESN specification. The hyperparameter sweep involved 4,752,000 ESN fits on the *Parameter* dataset. However, a total runtime for this tuning phase was not recorded. The runtimes reported in Tables 6 and 7 therefore refer only to fitting and forecasting with the already selected ESN configuration on the *Forecast* dataset. They should be interpreted as deployment runtimes rather than as the total computational cost of ESN tuning and forecasting. Computational runtimes are not reported for MLP, RNN, and ES-RNN because their forecasts were taken from the published M4 results. The reported runtimes should be interpreted as indicative values because execution times may vary between runs due to fluctuations in system load and background processes.

References

Akaike, H. (1974). A new look at the statistical model identification. *IEEE Transactions on Automatic Control*, 19(6):716–723.

Assimakopoulos, V. and Nikolopoulos, K. (2000). The theta model: A decomposition approach to forecasting. *International Journal of Forecasting*, 16:521–530.

Babinec, v. and Pospíšchal, J. (2012). Modular echo state neural networks in time series prediction. *Computing and Informatics*, 30(2):321–334.

Bergstra, J. and Bengio, Y. (2012). Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13(10):281–305.

Chouikhi, N., Ammar, B., Rokbani, N., and Alimi, A. M. (2017). Pso-based analysis of echo state network parameters for time series forecasting. *Applied Soft Computing*, 55:211–225.

Cleveland, R. B., Cleveland, W. S., McRae, J. E., and Terpenning, I. (1990). Stl: A seasonal-trend decomposition procedure based on loess. *Journal of Official Statistics*, 6(1):3–73.

Gallicchio, C. and Micheli, A. (2017). Deep echo state network (DeepESN): A brief survey. *arXiv preprint arXiv:1712.04323*.

Gallicchio, C., Micheli, A., and Pedrelli, L. (2018). Design of deep echo state networks. *Neural Networks*, 108:33–47.

Hannan, E. J. and Quinn, B. G. (1979). The determination of the order of an autoregression. *Journal of the Royal Statistical Society. Series B (Methodological)*, 41(2):190–195.

Hastie, T., Tibshirani, R., and Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer, New York, 2nd edition.

Hoerl, A. E. and Kennard, R. W. (1970). Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67.

Holm, S. (1979). A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70.

Hurvich, C. M. and Tsai, C.-L. (1989). Regression and time series model selection in small samples. *Biometrika*, 76(2):297–307.

Hyndman, R., Athanasopoulos, G., Bergmeir, C., Caceres, G., Chhay, L., O'Hara-Wild, M., Petropoulos, F., Razbash, S., Wang, E., and Yasmeen, F. (2026). *forecast: Forecasting functions for time series and linear models*. R package version 9.0.1.

Hyndman, R., Koehler, A., Ord, K., and Snyder, R. (2008). *Forecasting with exponential smoothing. The state space approach*.

Hyndman, R. J. and Athanasopoulos, G. (2021). *Forecasting: principles and practice*. OTexts: Melbourne, Australia, 3rd edition.

Hyndman, R. J. and Khandakar, Y. (2008). Automatic time series forecasting: The forecast package for r. *Journal of Statistical Software, Articles*, 27(3):1–22.

Hyndman, R. J. and Koehler, A. B. (2006). Another look at measures of forecast accuracy. *International Journal of Forecasting*, 22(4):679–688.

Häußer, A. (2026a). *echos: Echo State Networks for Time Series Modeling and Forecasting*. R package version 1.0.4.

Häußer, A. (2026b). *tscv: Functions and Utilities for Tidy Time Series Forecasting and Time Series Cross-Validation*. R package version 1.0.0.

Jaeger, H. (2001). The "echo state" approach to analysing and training recurrent neural networks-with an erratum note. *Bonn, Germany: German National Research Center for Information Technology GMD Technical Report*, 148(34):13.

Jaeger, H. (2002). Tutorial on training recurrent neural networks, covering bppt, rtrl, ekf and the echo state network approach. *GMD-Forschungszentrum Informationstechnik, 2002.*, 5.

Kwiatkowski, D., Phillips, P. C., Schmidt, P., and Shin, Y. (1992). Testing the null hypothesis of stationarity against the alternative of a unit root: How sure are we that economic time series have a unit root? *Journal of Econometrics*, 54(1-3):159–178.

Livera, A. M. D., Hyndman, R. J., and Snyder, R. D. (2011). Forecasting time series with complex seasonal patterns using exponential smoothing. *Journal of the American Statistical Association*, 106(496):1513–1527.

Lukoševičius, M. (2012). A practical guide to applying echo state networks. In *Neural Networks: Tricks of the Trade: Second Edition*, pages 659–686. Springer.

Lukoševičius, M. and Jaeger, H. (2009). Reservoir computing approaches to recurrent neural network training. *Computer Science Review*, 3(3):127–149.

Maat, J. R., Gianniotis, N., and Protopapas, P. (2018). Efficient optimization of echo state networks for time series datasets. In *2018 International Joint Conference on Neural Networks (IJCNN)*, pages 1–7.

Makridakis, S. and Hibon, M. (2000). The m3-competition: results, conclusions and implications. *International Journal of Forecasting*, 16(4):451–476. The M3-Competition.

Makridakis, S., Spiliotis, E., and Assimakopoulos, V. (2020). The M4 competition: 100,000 time series and 61 forecasting methods. *International Journal of Forecasting*, 36(1):54–74. M4 Competition.

Makridakis, S. G., Wheelwright, S. C., and Hyndman, R. J. (1998). *Forecasting : methods and applications*. John Wiley & Sons, New York, 3rd edition.

O'Hara-Wild, M., Hyndman, R., and Wang, E. (2026a). *fable: Forecasting Models for Tidy Time Series*. R package version 0.5.0.

O'Hara-Wild, M., Hyndman, R., and Wang, E. (2026b). *feasts: Feature Extraction and Statistics for Time Series*. R package version 0.5.0.

R Core Team (2025). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.

Schwarz, G. (1978). Estimating the dimension of a model. *Annals of Statistics*, 6(2):461–464.

Smyl, S. (2020). A hybrid method of exponential smoothing and recurrent neural networks for time series forecasting. *International Journal of Forecasting*, 36(1):75–85.

Spiliotis, E., Kouloumos, A., Assimakopoulos, V., and Makridakis, S. (2020). Are forecasting competitions data representative of the reality? *International Journal of Forecasting*, 36(1):37–53. M4 Competition.

Sugiura, N. (1978). Further analysis of the data by akaike's information criterion and the finite corrections. *Communications in Statistics - Theory and Methods*, 7(1):13–26.

Viehweg, J., Teutsch, P., and Mäder, P. (2025). A systematic study of echo state networks topologies for chaotic time series prediction. *Neurocomputing*, 618:129032.

Wilcoxon, F. (1945). Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83.

Xu, M. and Han, M. (2016). Adaptive elastic echo state network for multivariate time series prediction. *IEEE Transactions on Cybernetics*, 46(10):2173–2183.

Yao, X. and Wang, Z. (2019). Broad echo state network for multivariate time series prediction. *Journal of the Franklin Institute*, 356(9):4888–4906.